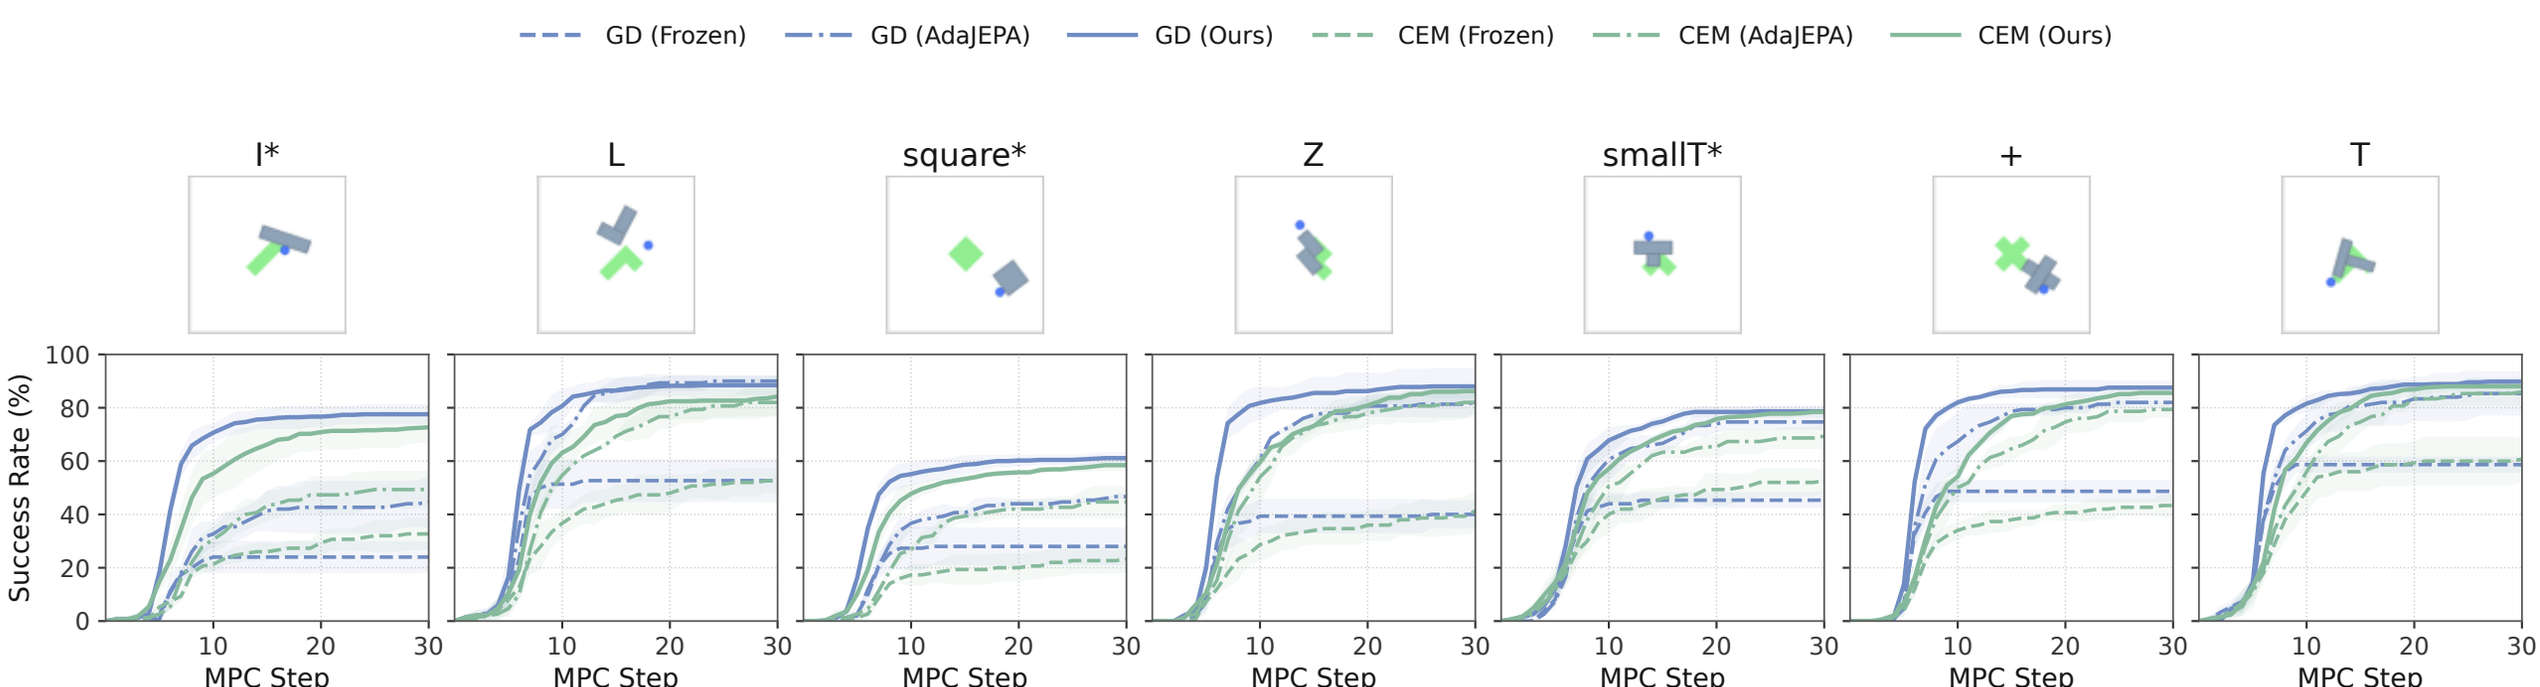
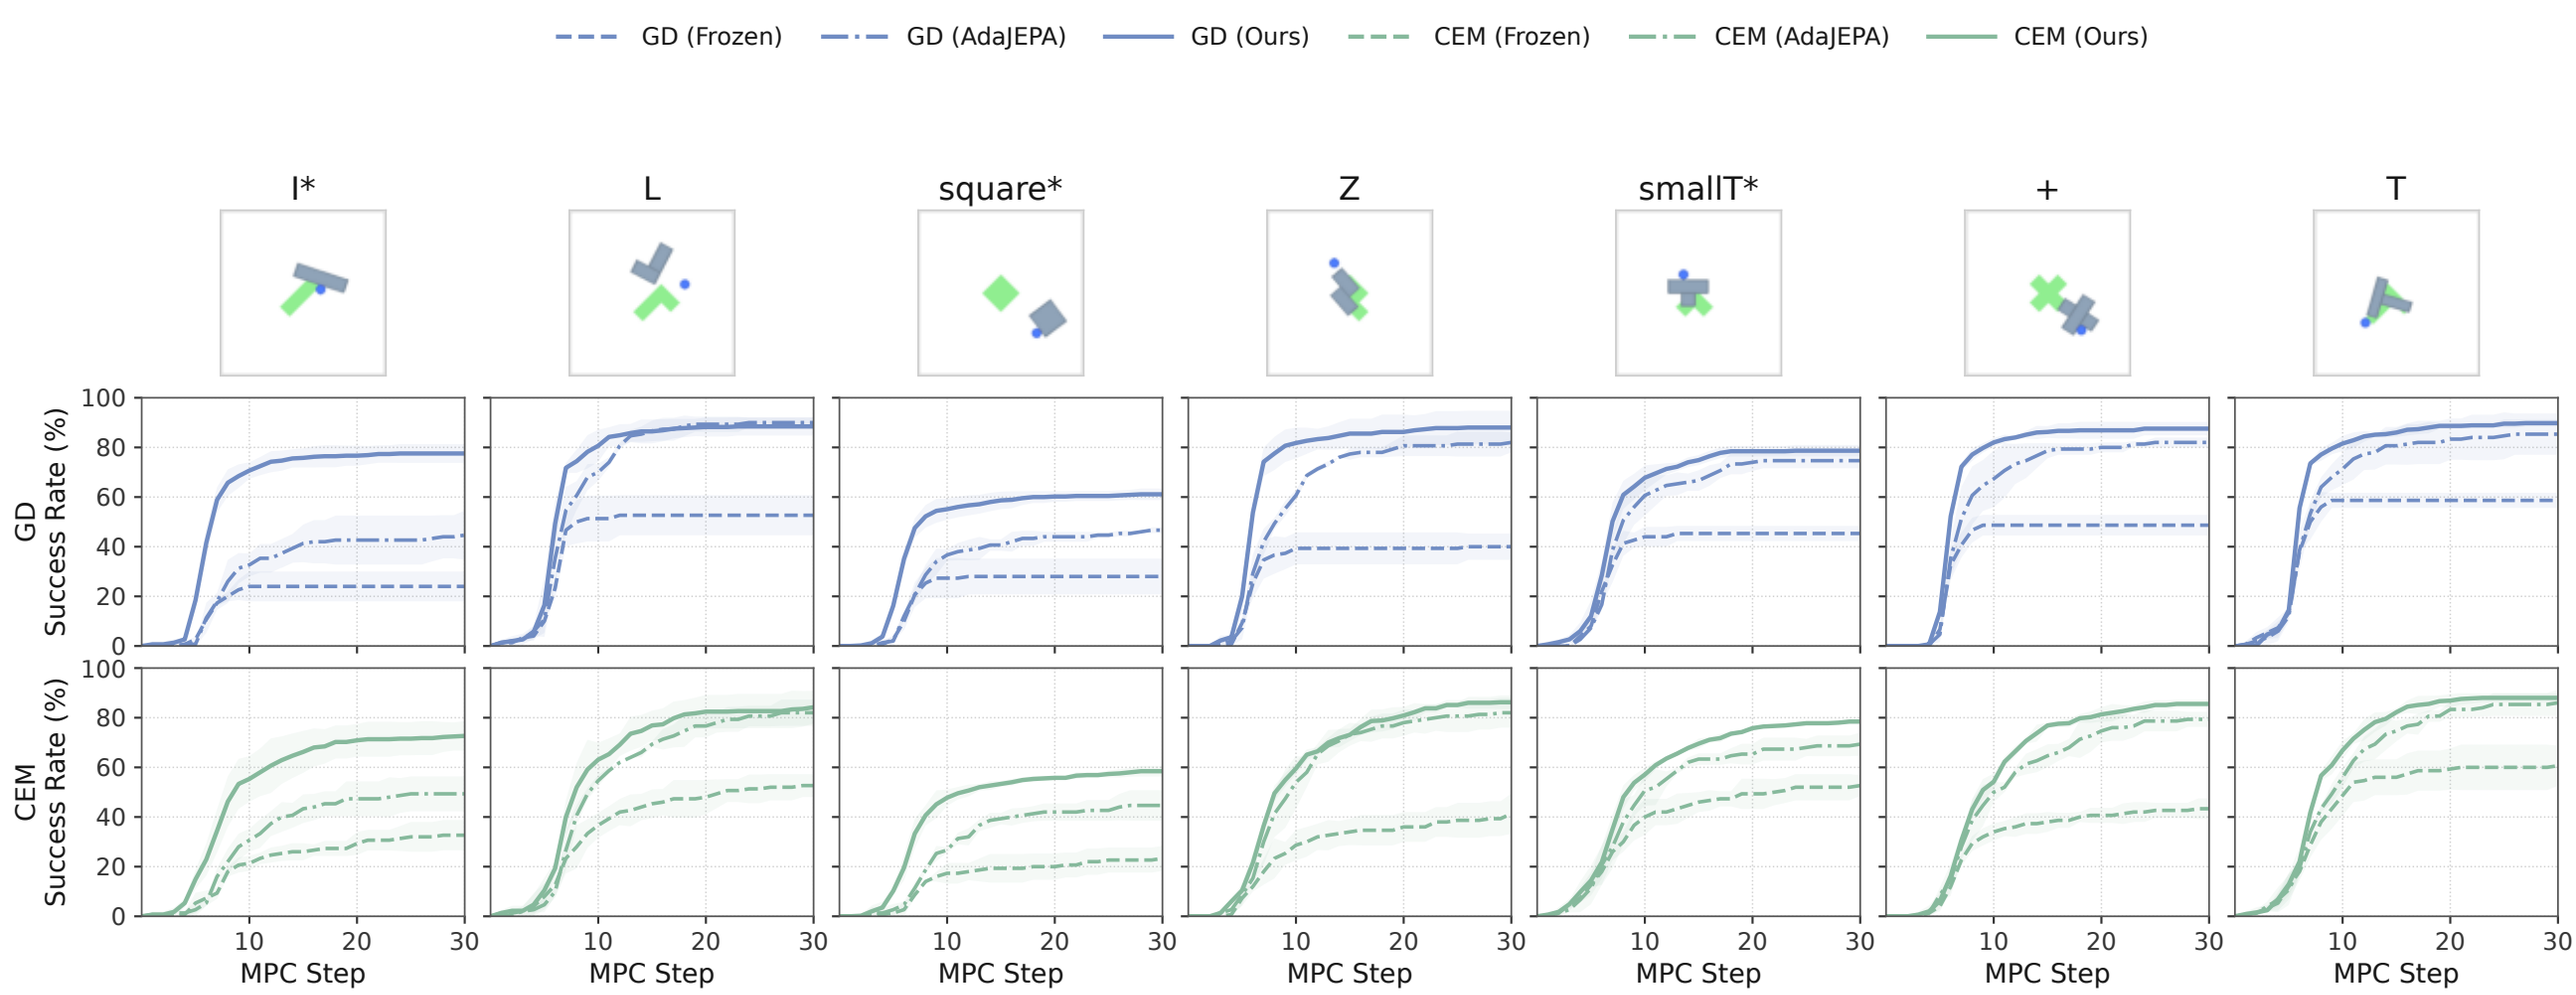
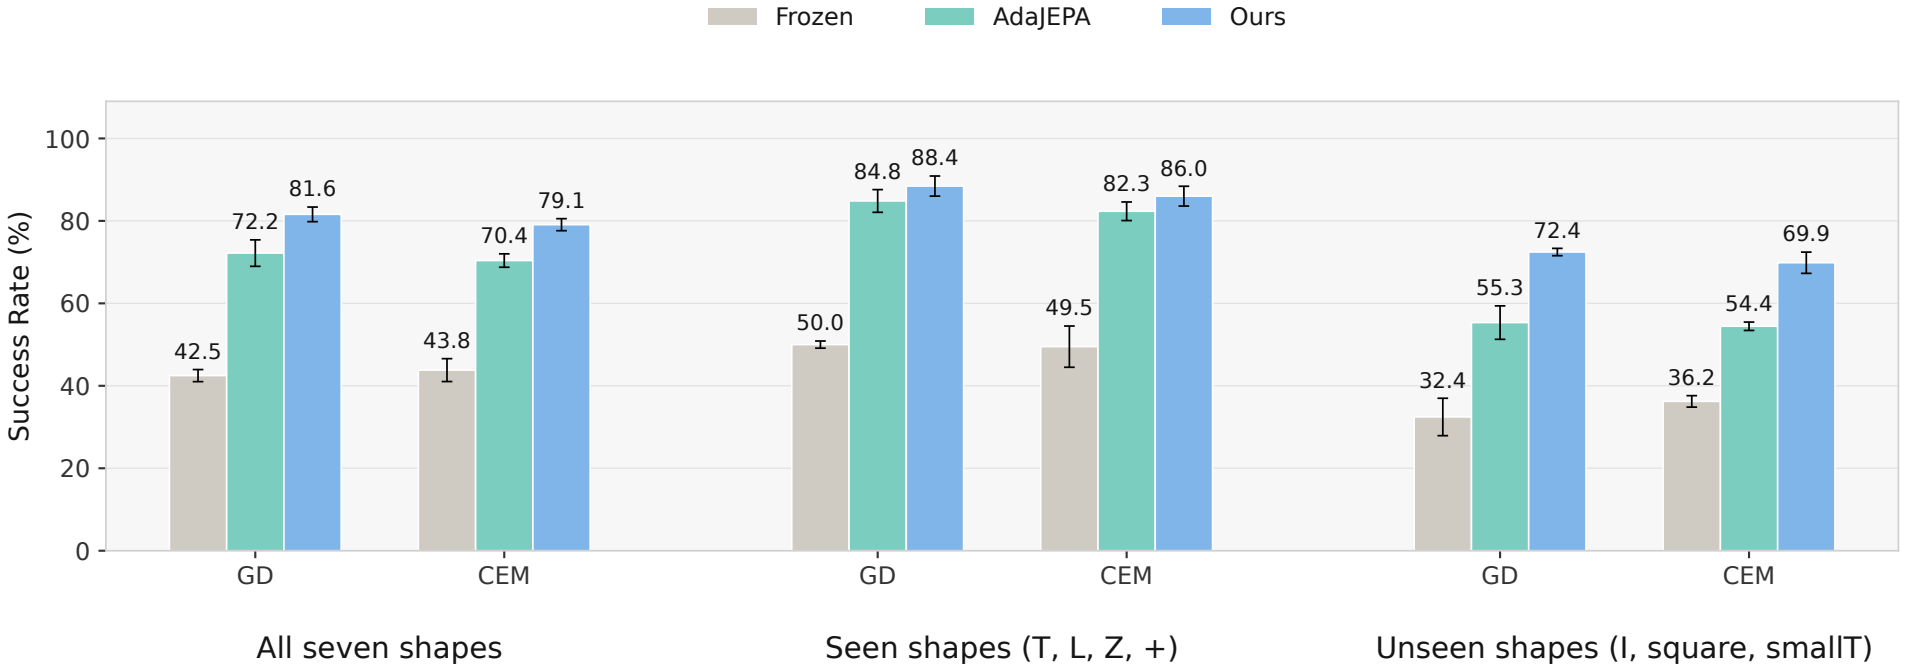


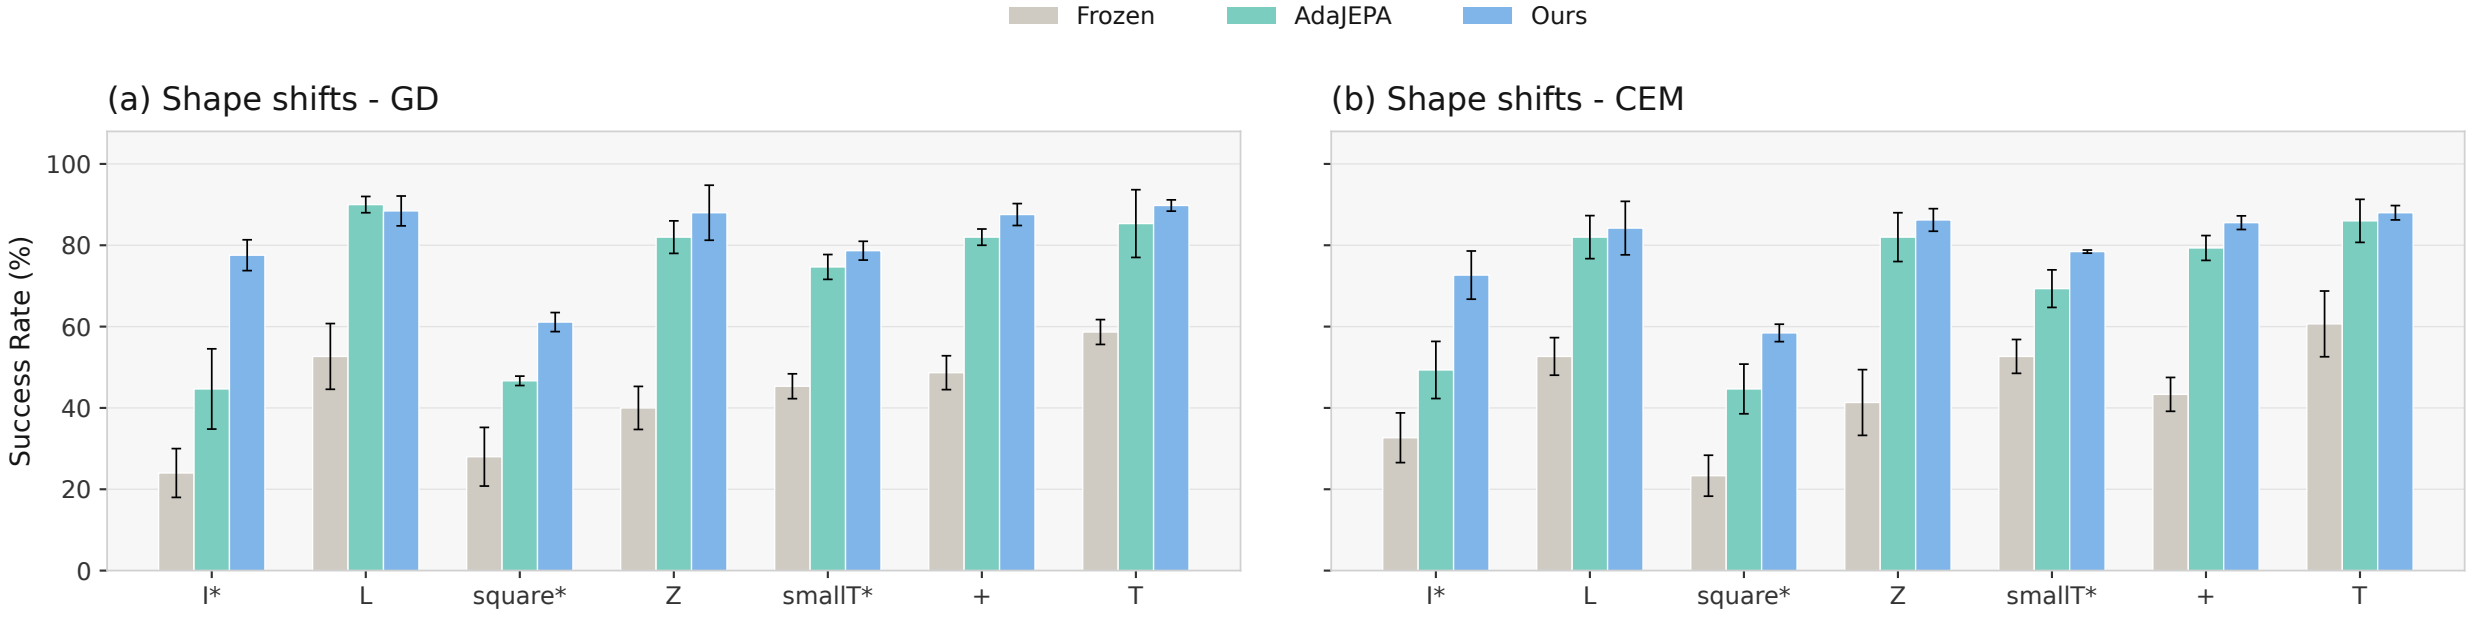
Mean +/- 1 SD across test seeds 2000/2100/2200. * shape not in our base model's training data. Ours: round-5 states of three chains, read-only; baselines: official shape-shift checkpoint.



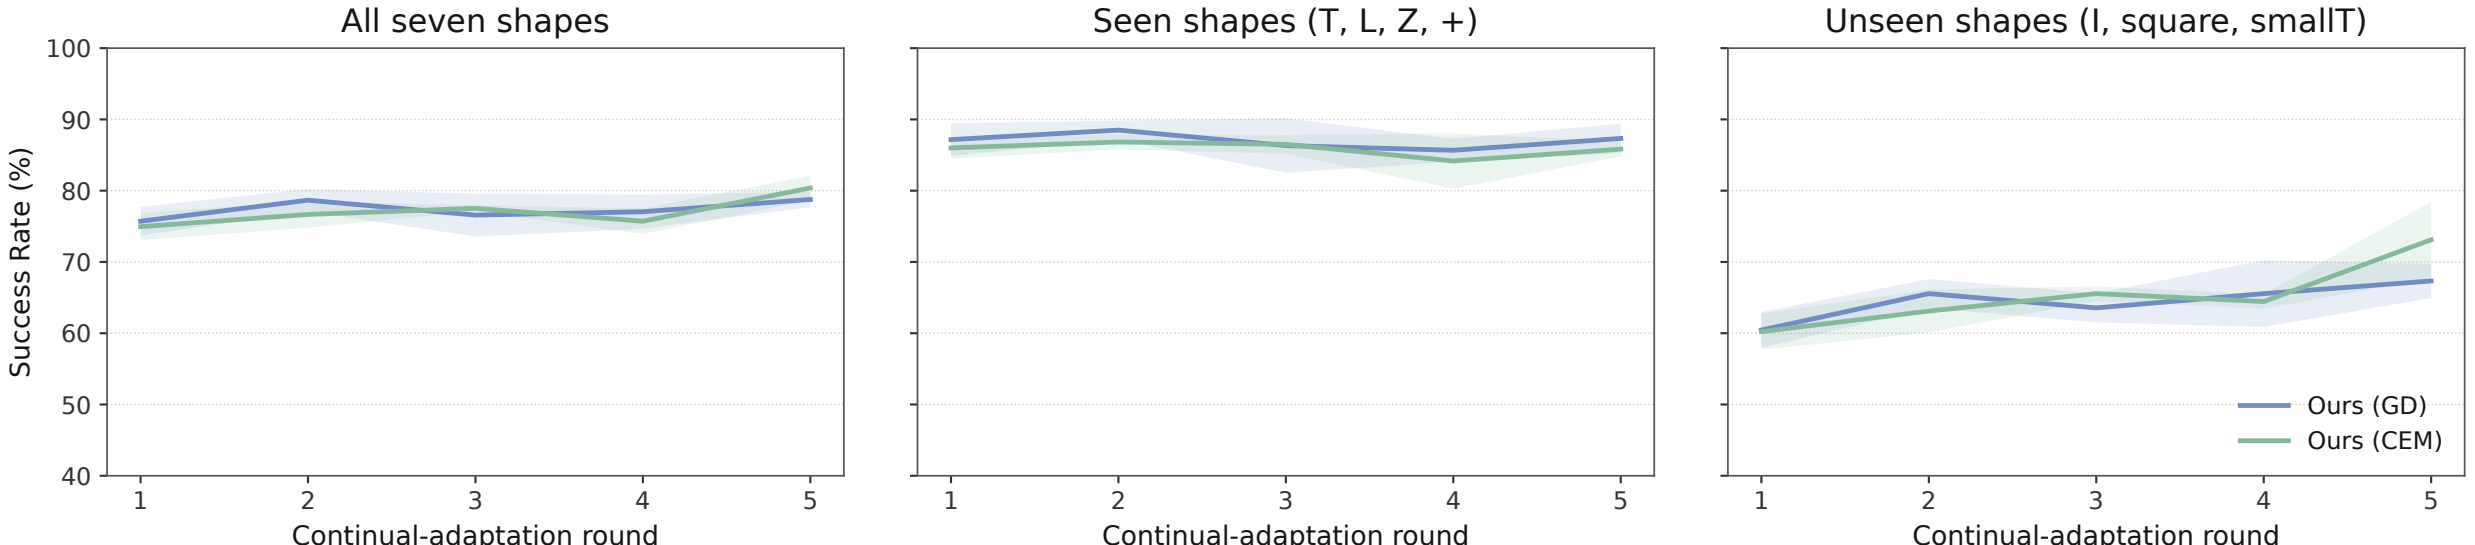
Mean +/- 1 SD across test seeds 2000/2100/2200. * shape not in our base model's training data. Ours: round-5 states of three chains, read-only; baselines: official shape-shift checkpoint.



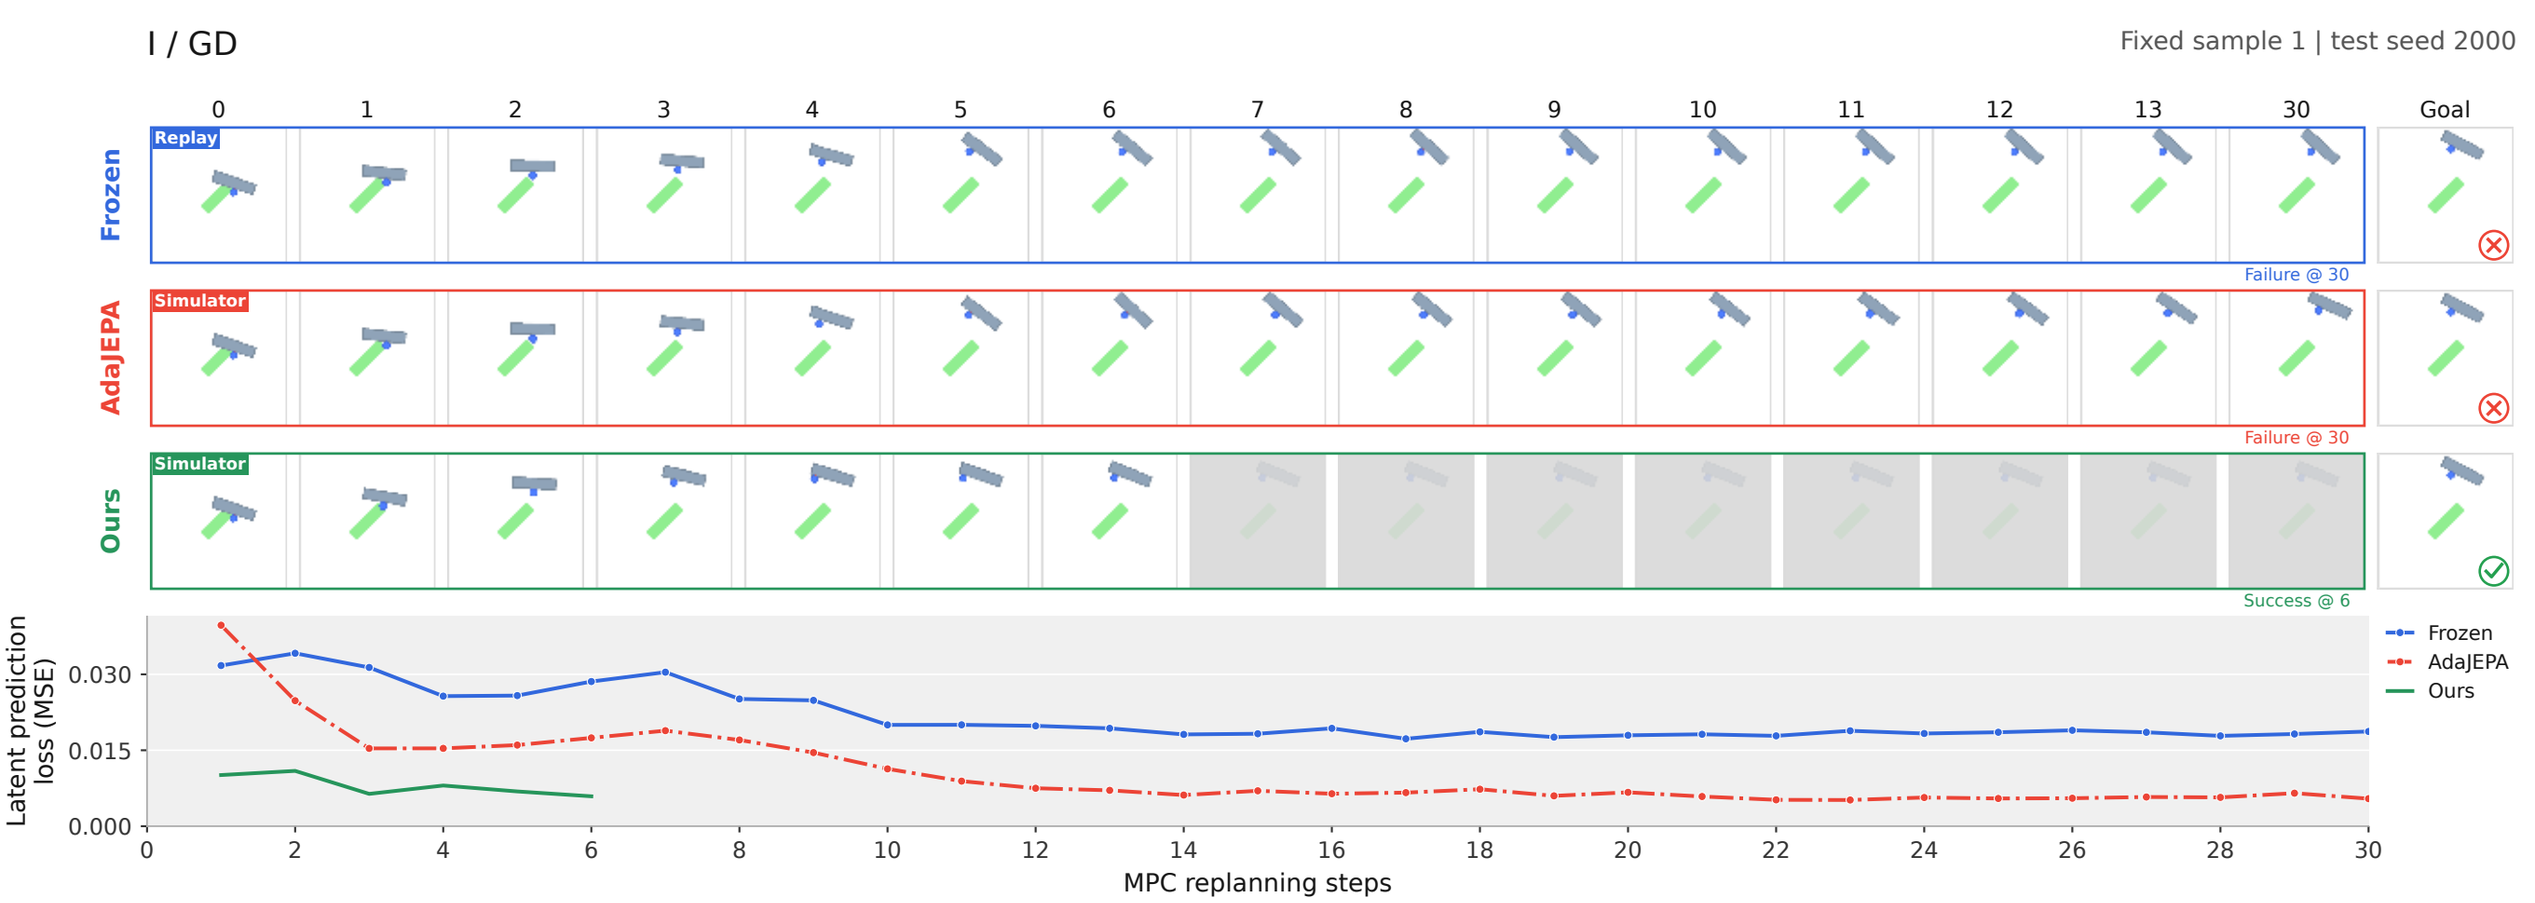
Error bars: +/- 1 SD across test seeds 2000/2100/2200. 1,050 episodes per method and planner (Ours: 3,150, three chains).



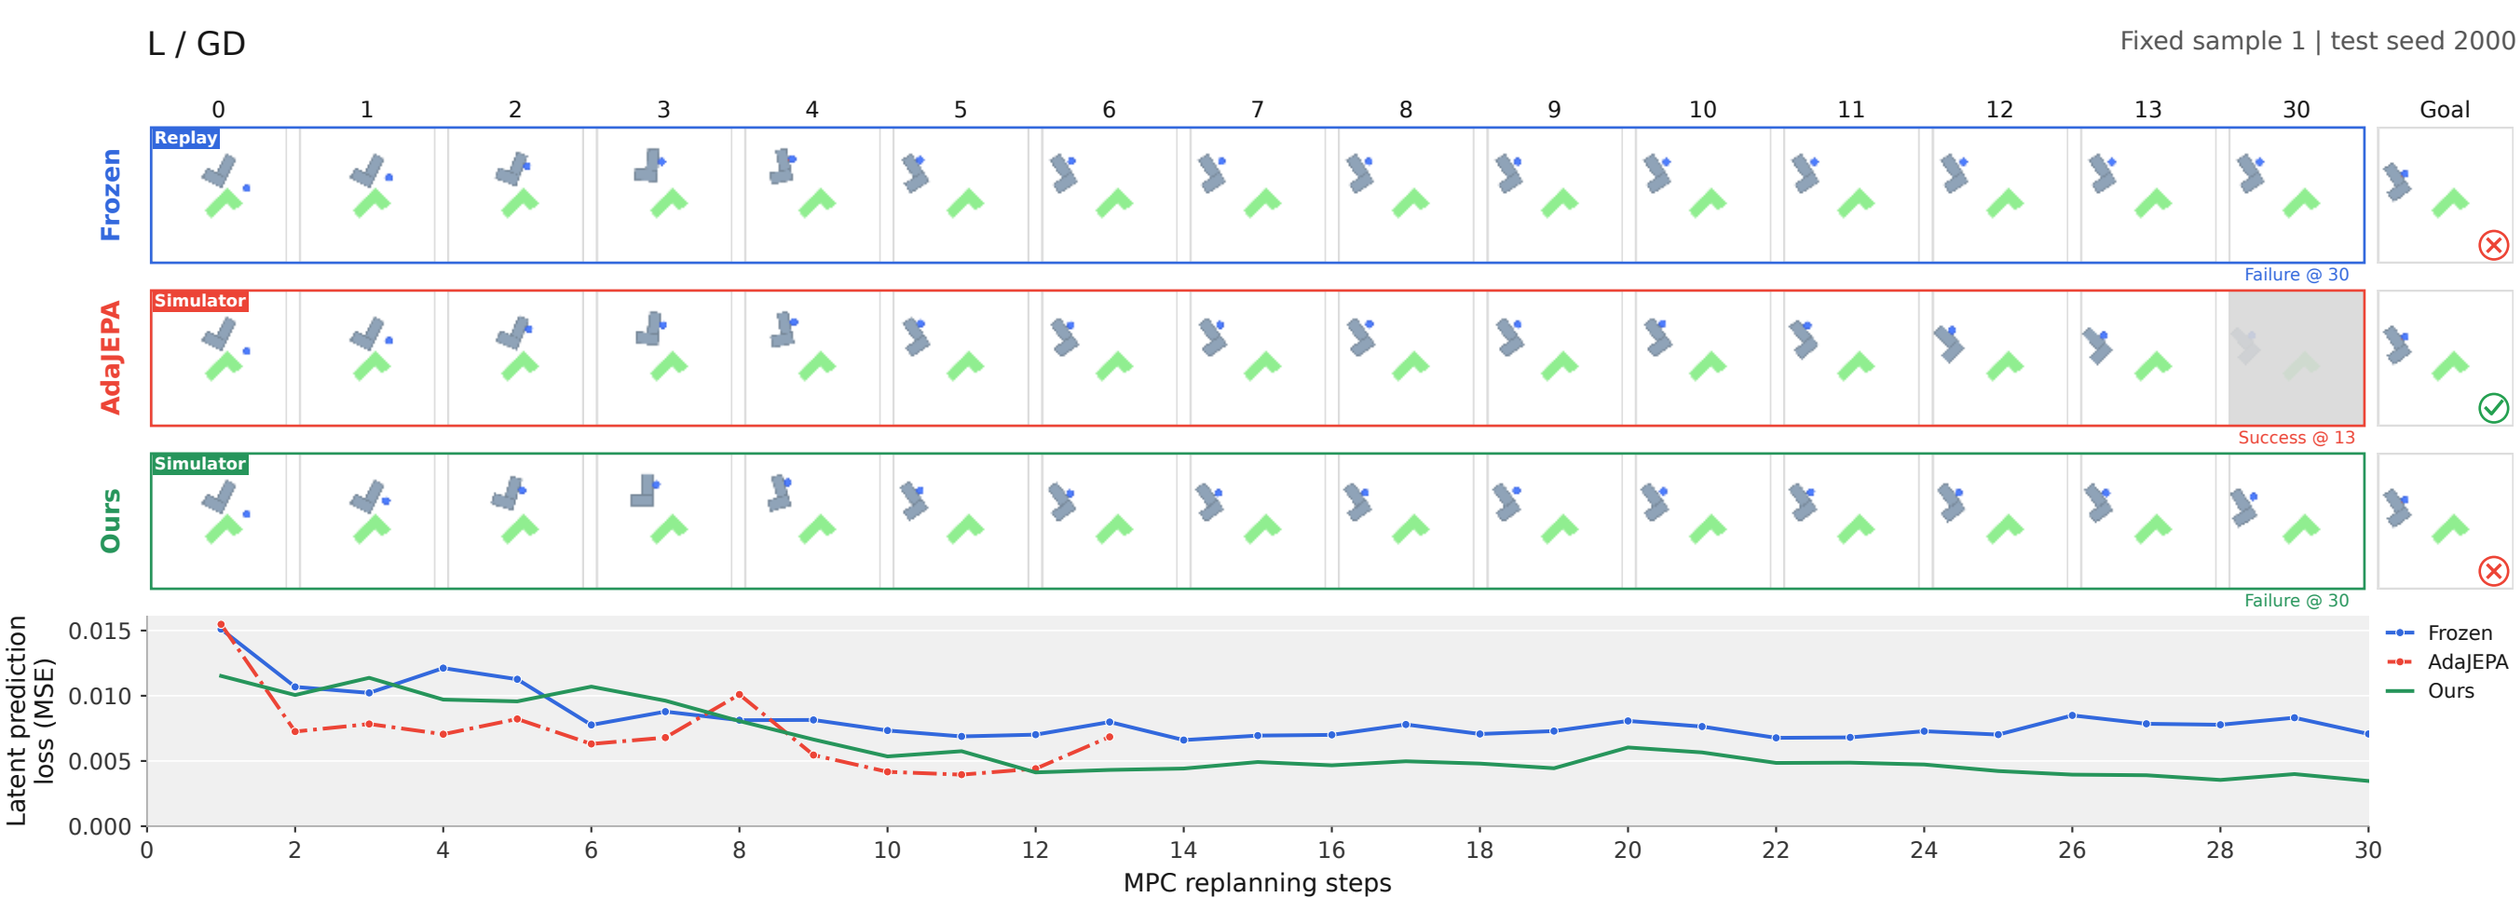
Mean +/- 1 SD across test seeds 2000/2100/2200; 150 episodes per shape per method (Ours: 450). * shape not in our base model's training data.



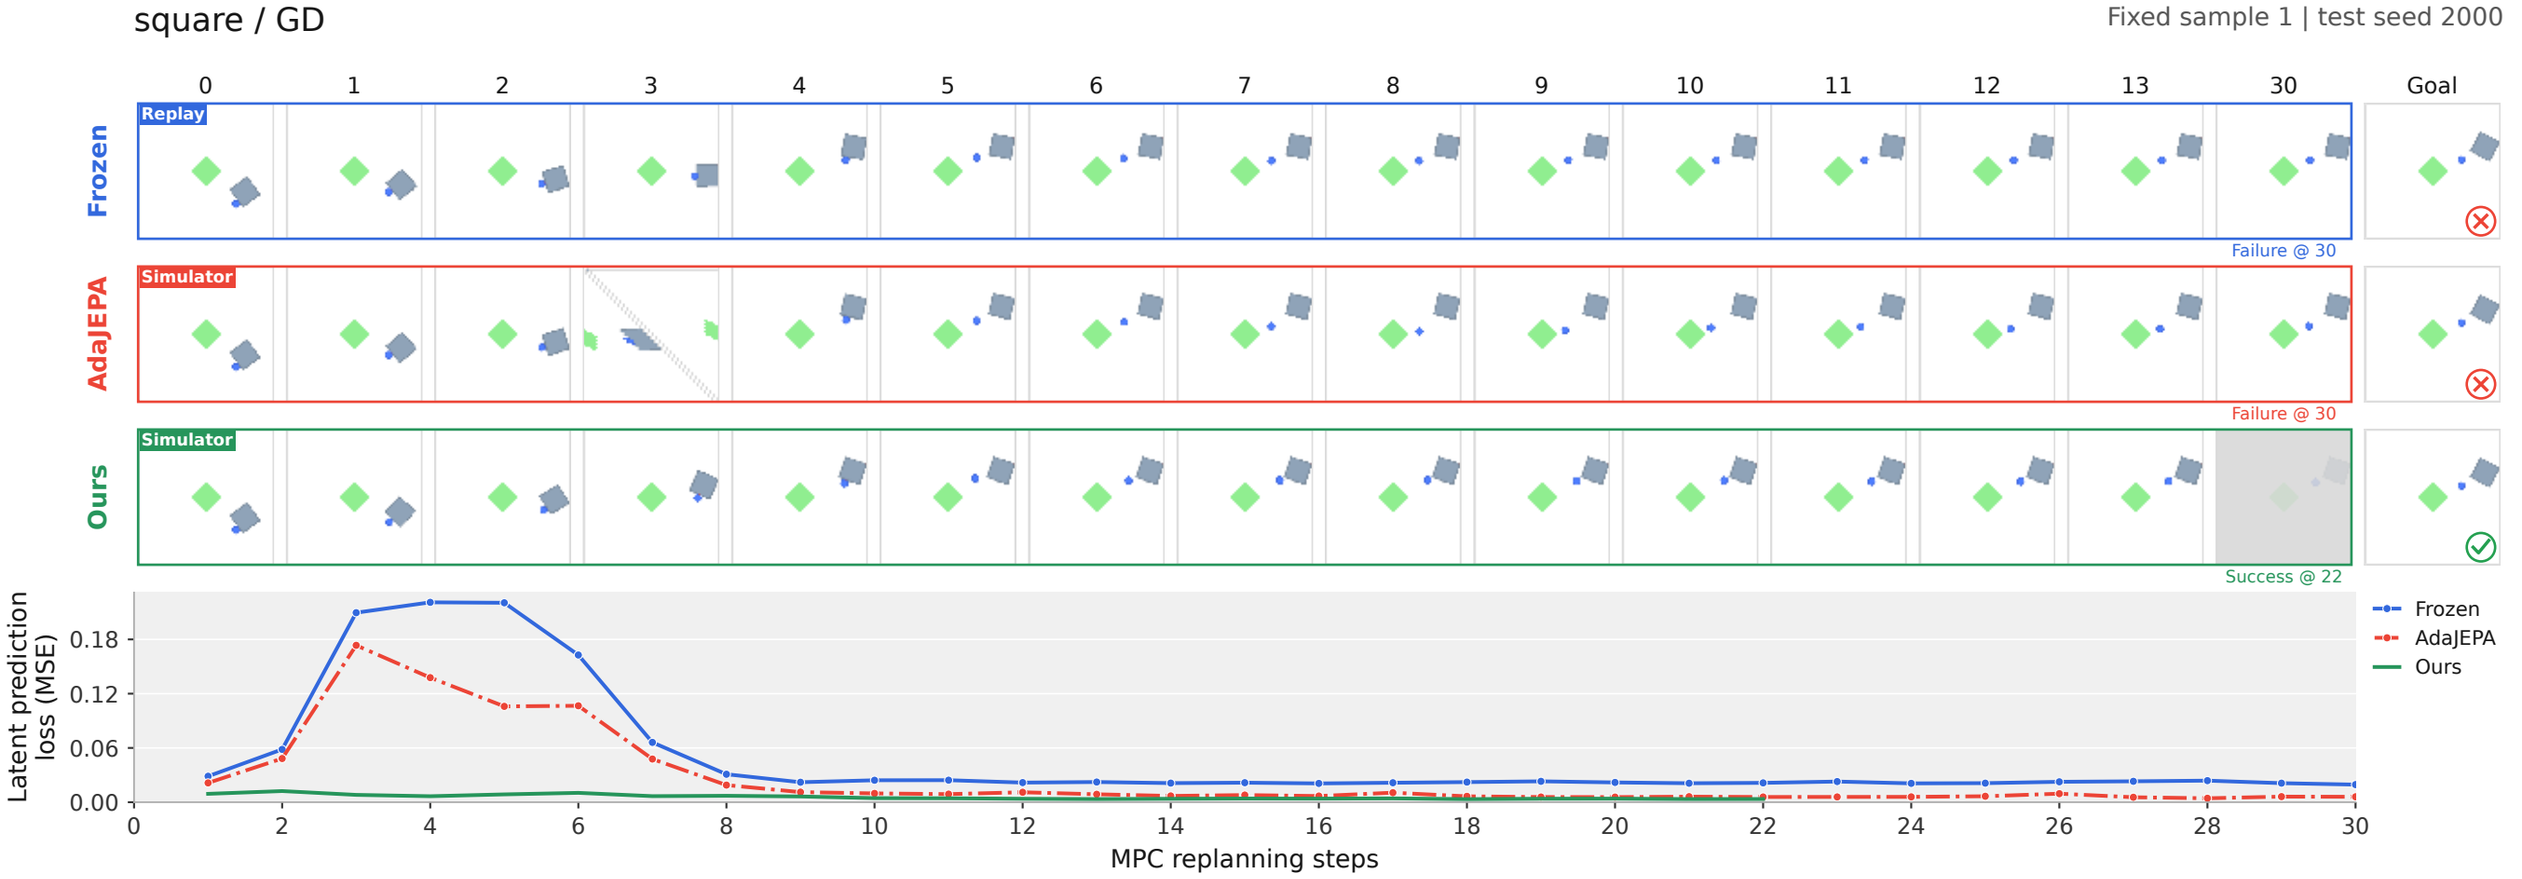
Mean +/- 1 SD across three chains (GD: streams S7t3, S7t3s400, S7t3s700; CEM: S7t3, S7t3b, S7t3s700). Each round: 7 shapes x 50 episodes with new test seeds; Ours keeps learning across rounds.



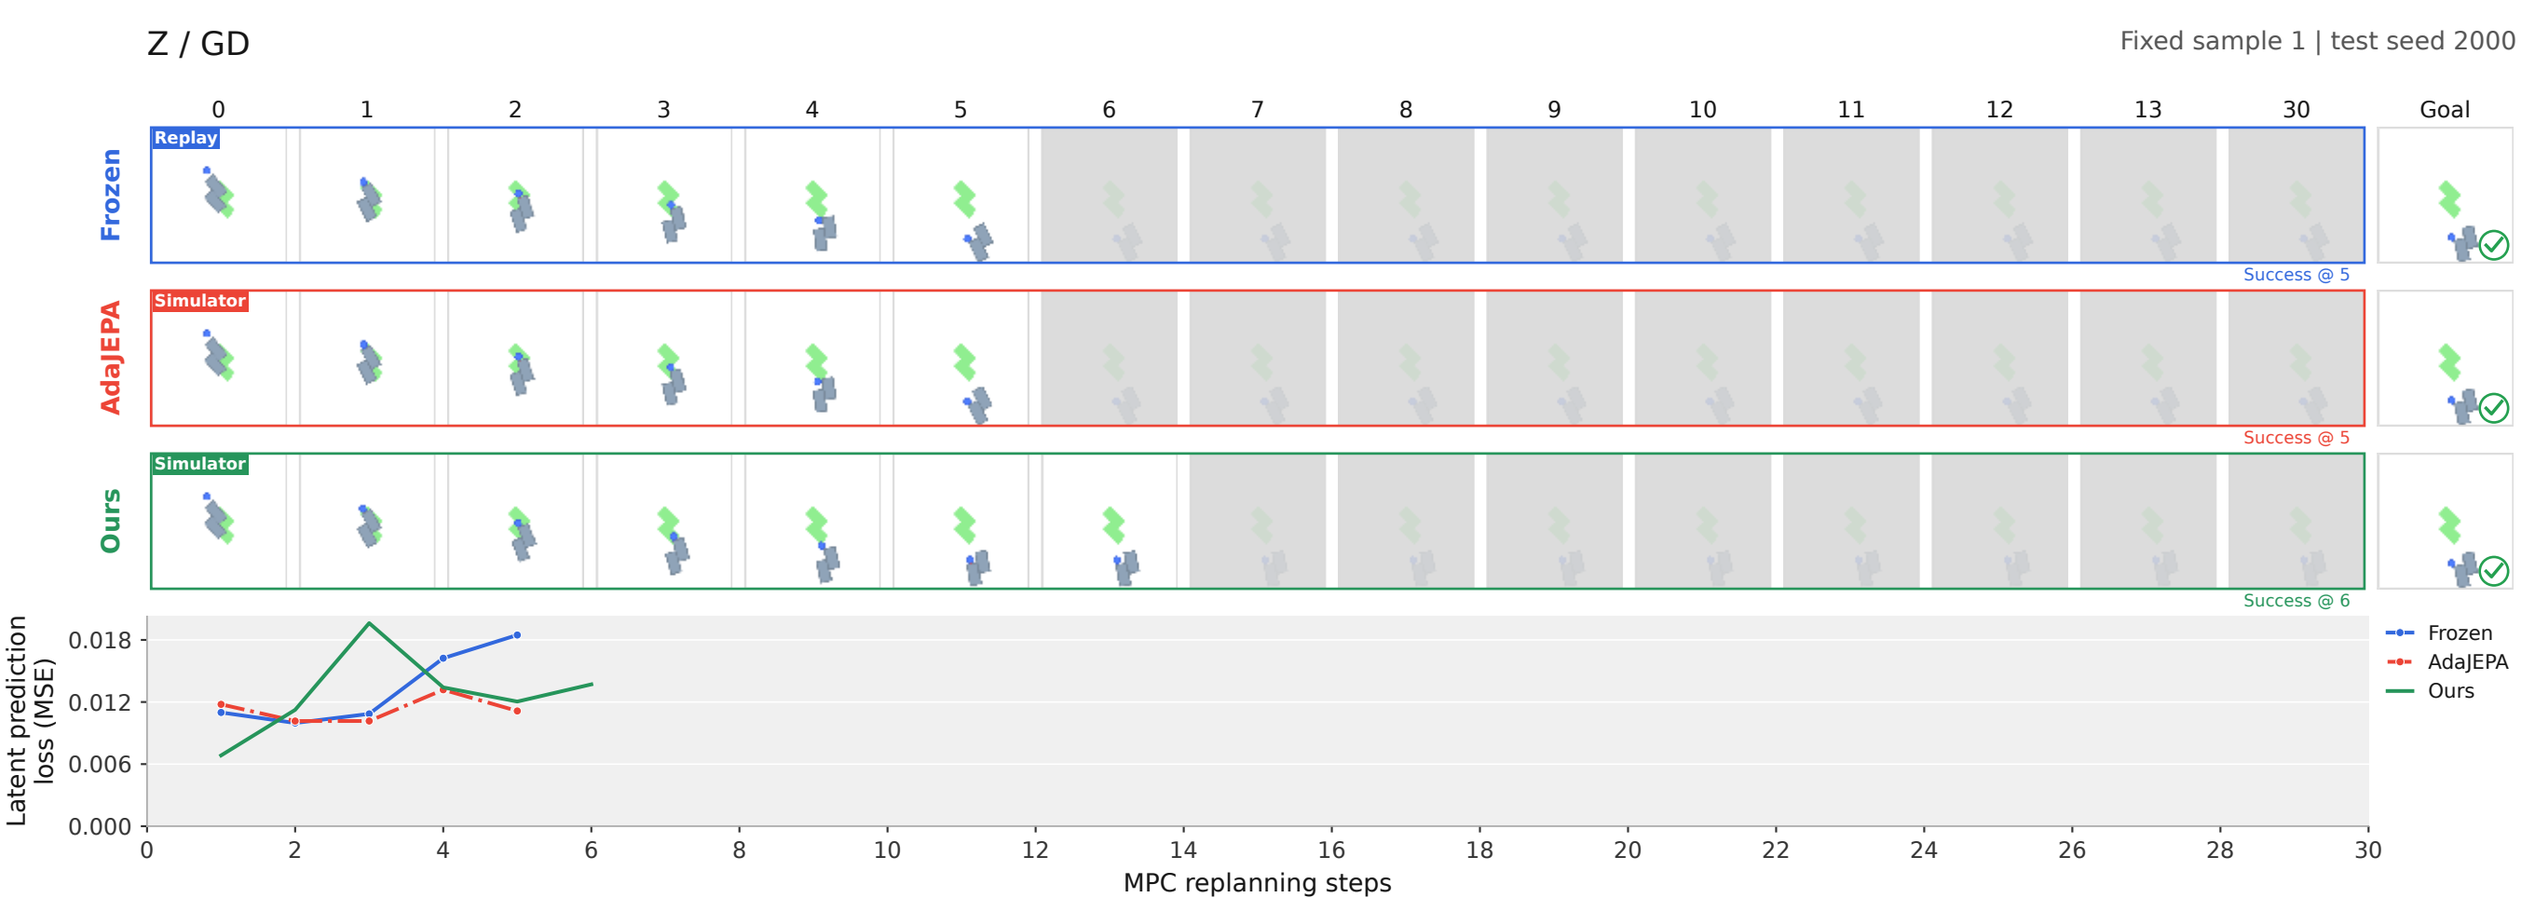
Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.



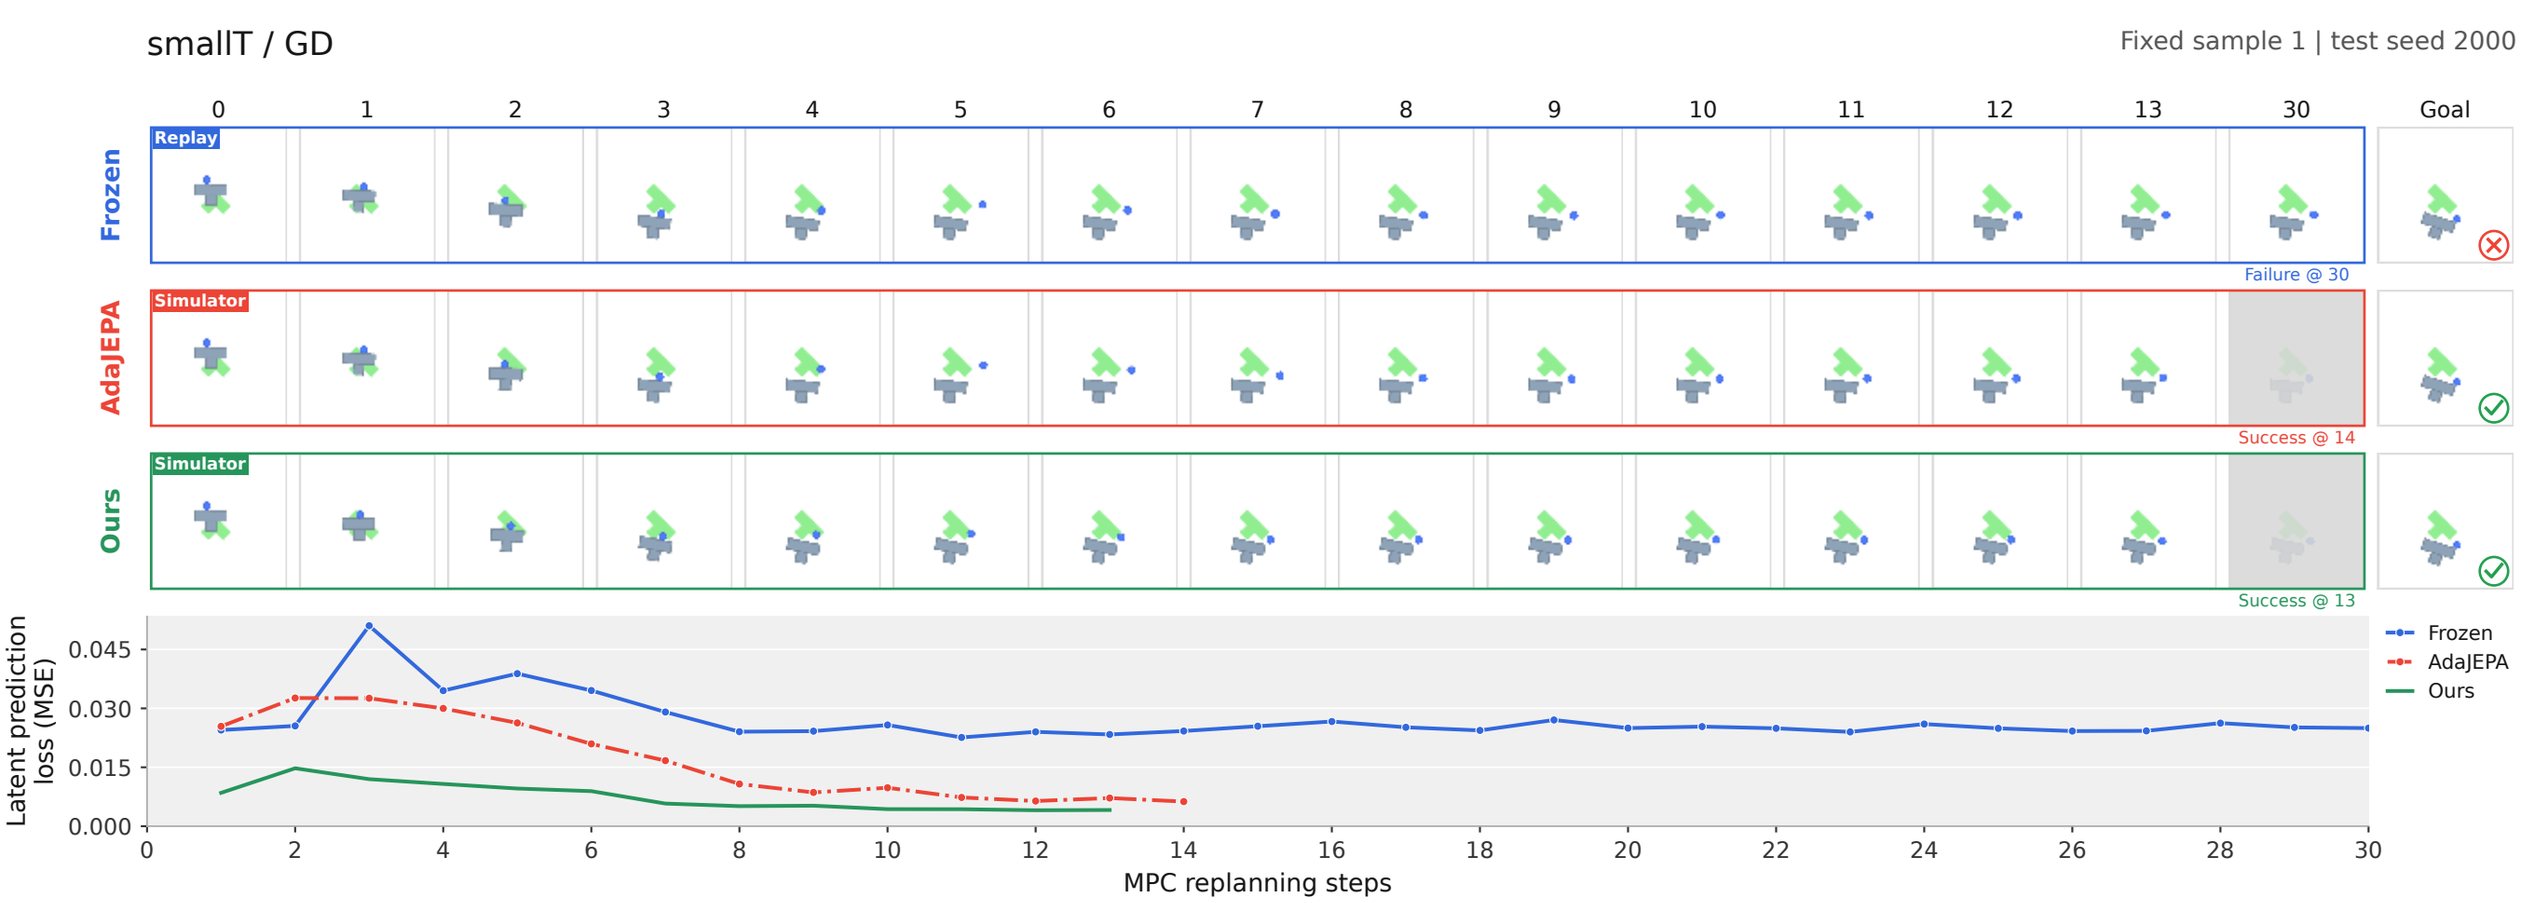
Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.



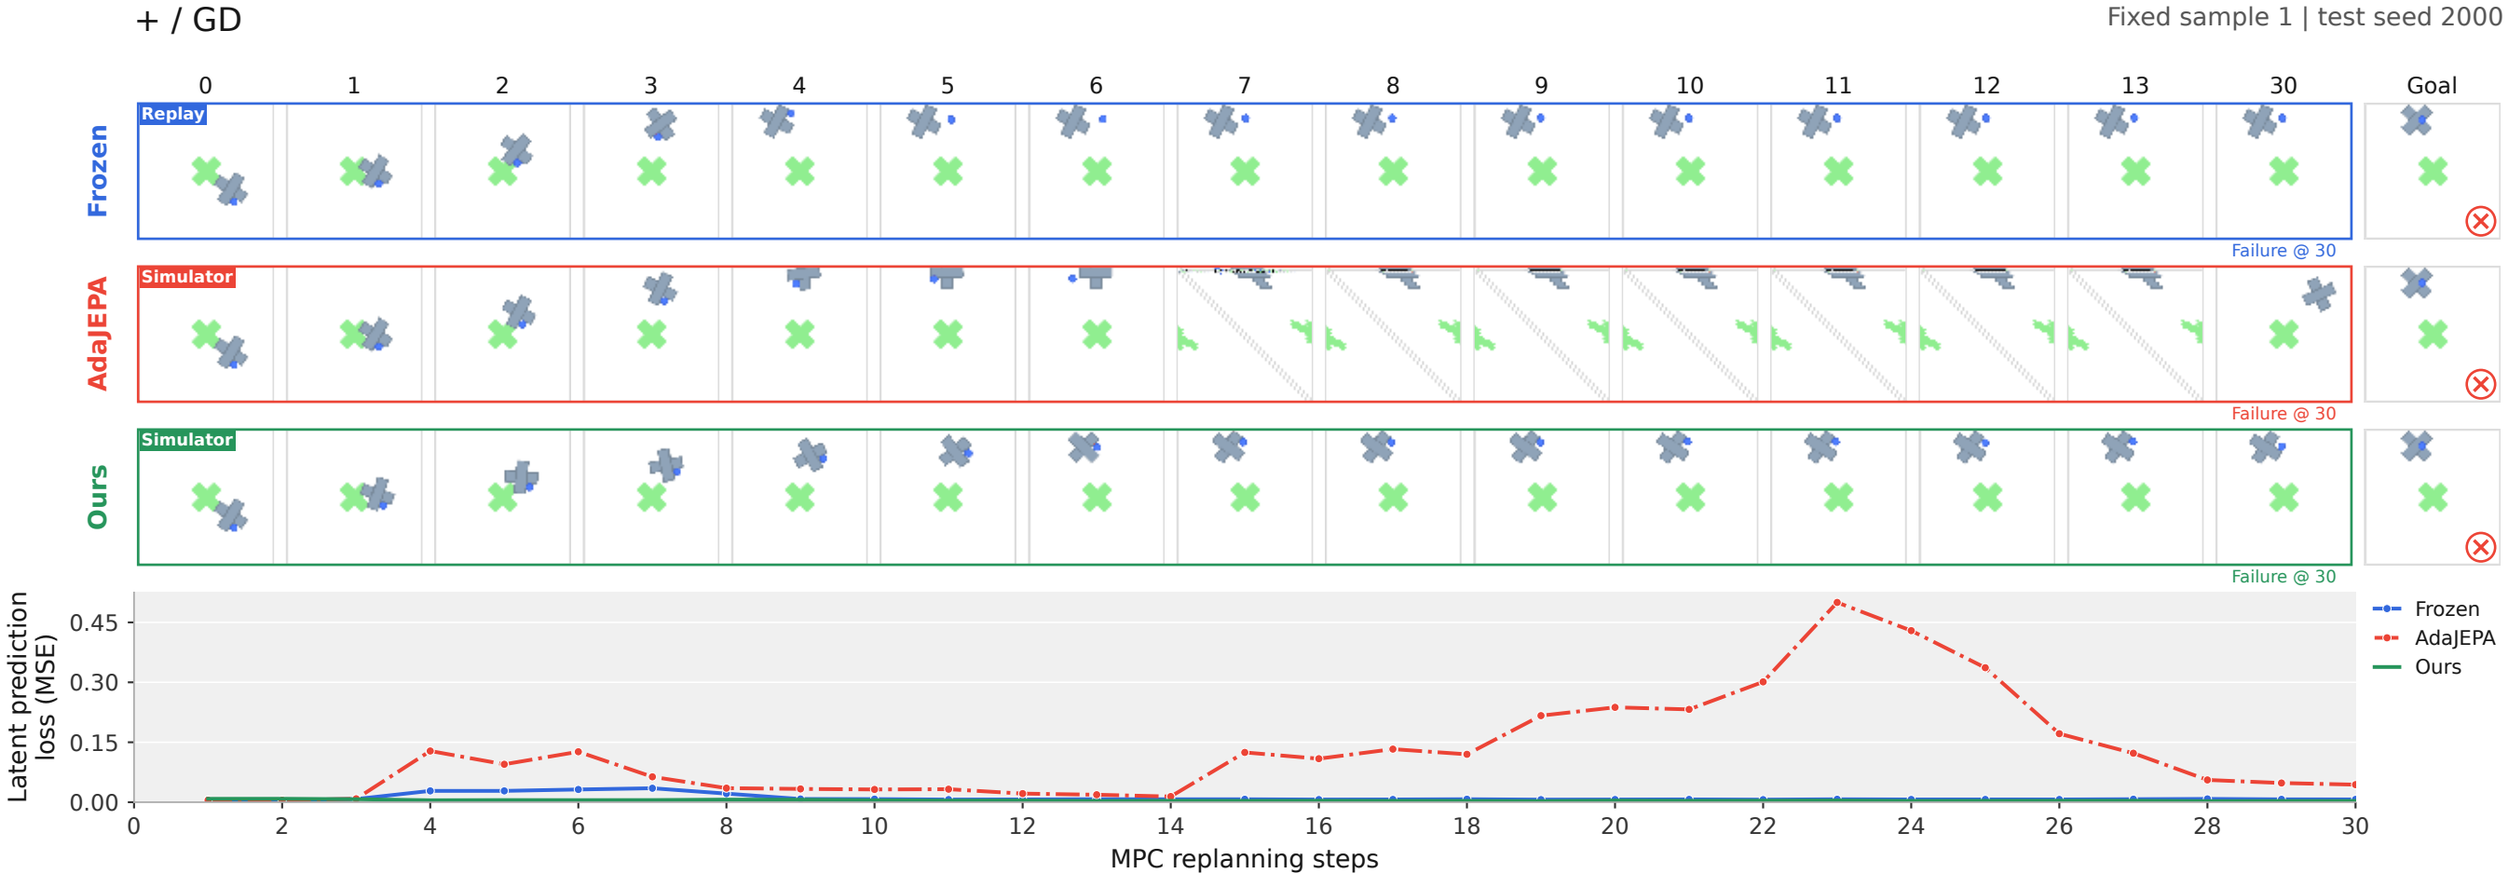
Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.



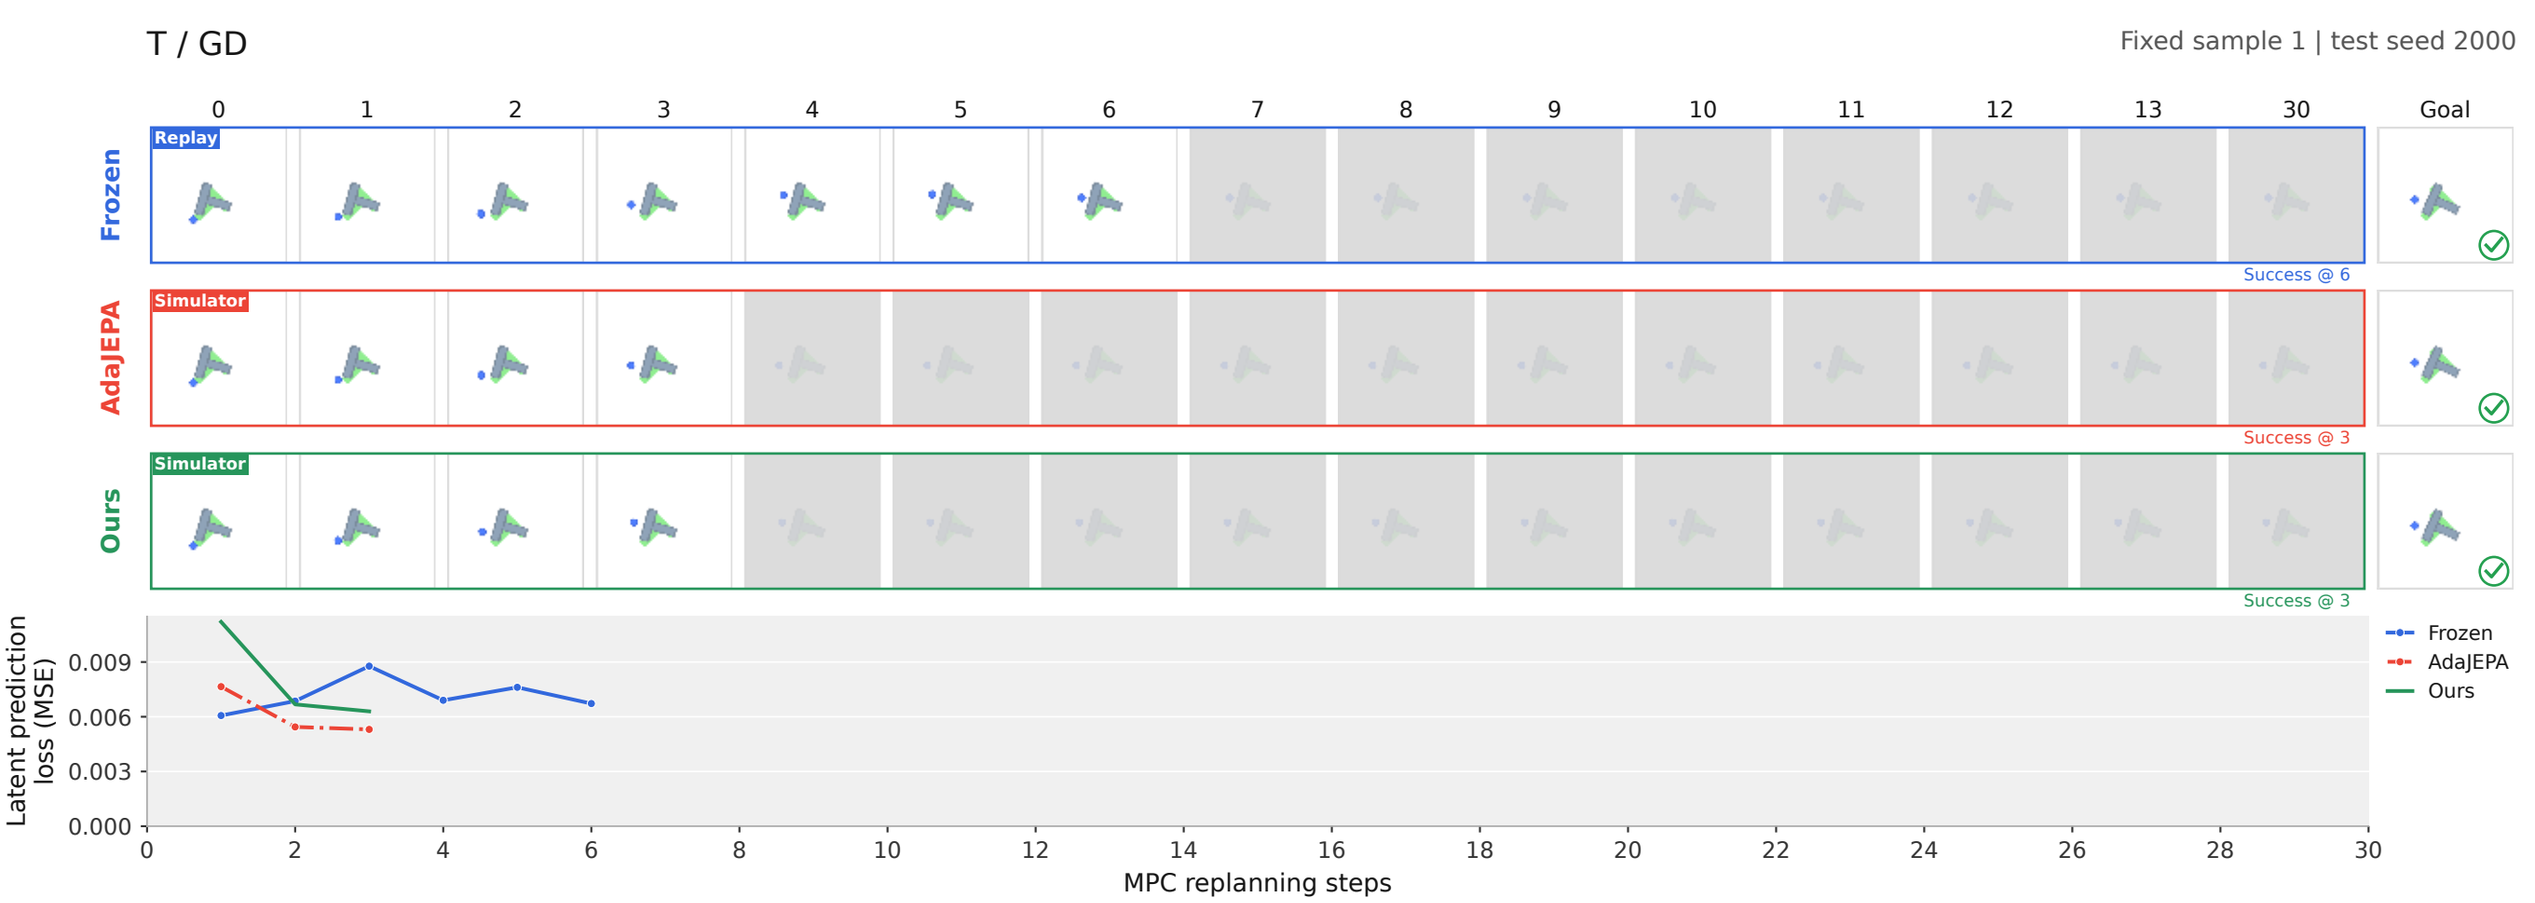
Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.



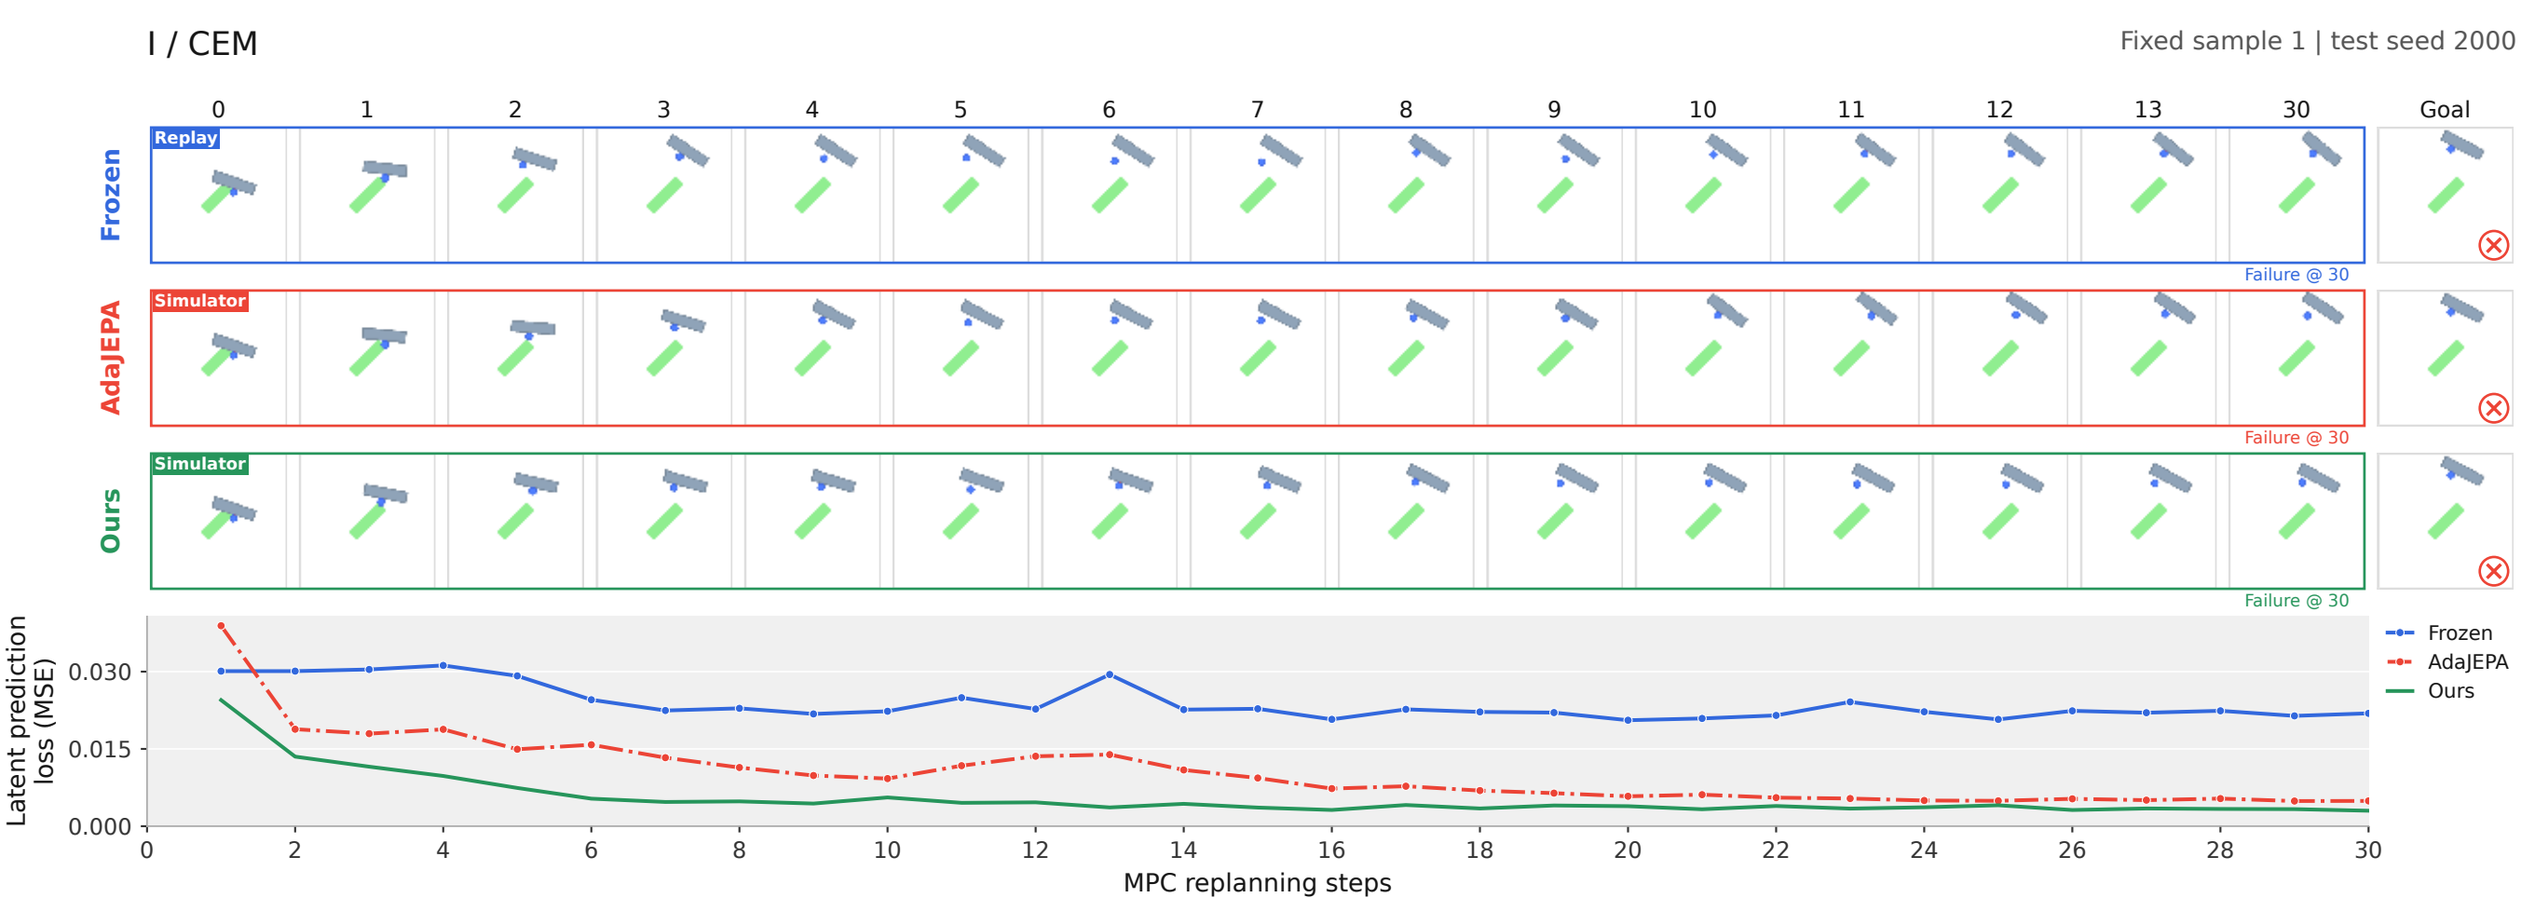
Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.



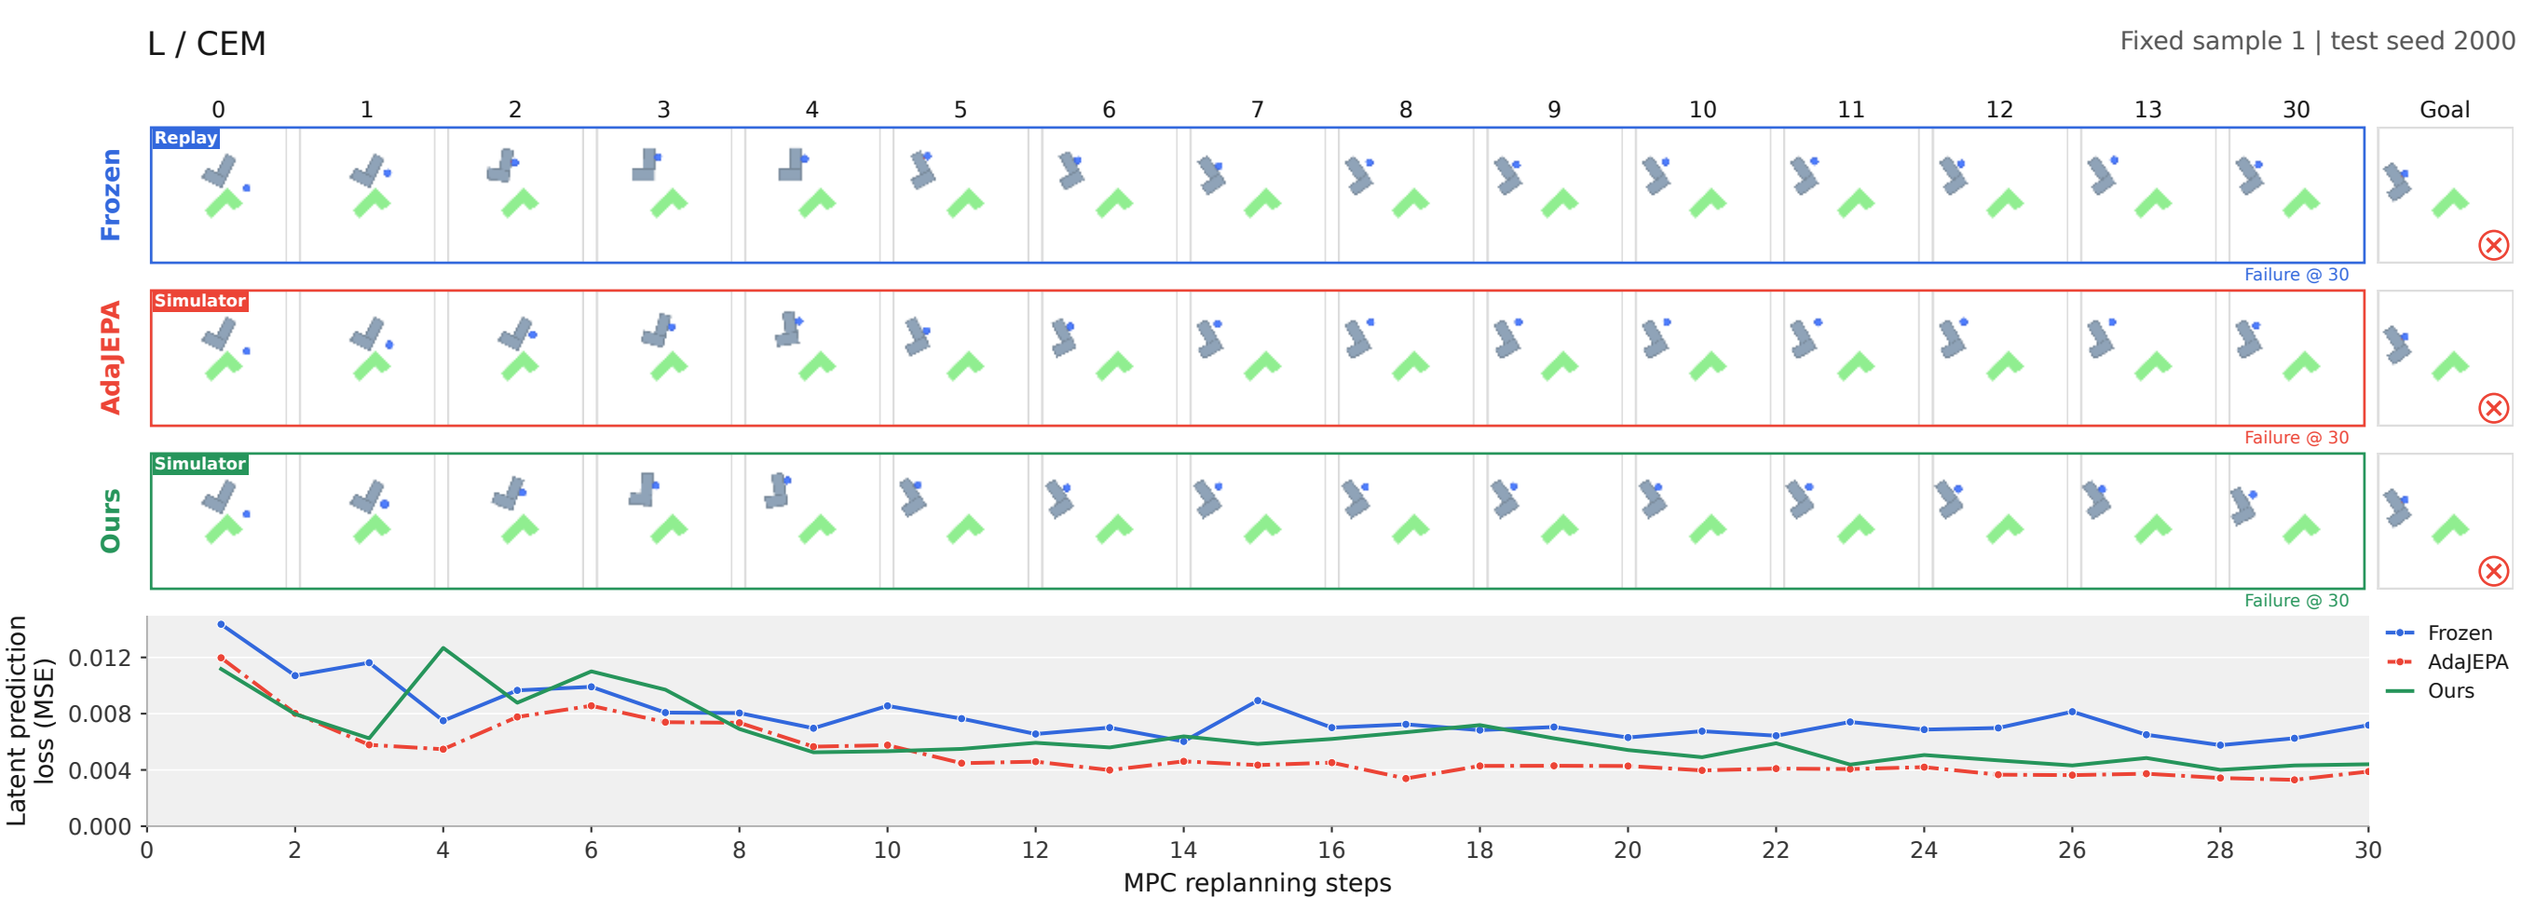
Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.



Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.

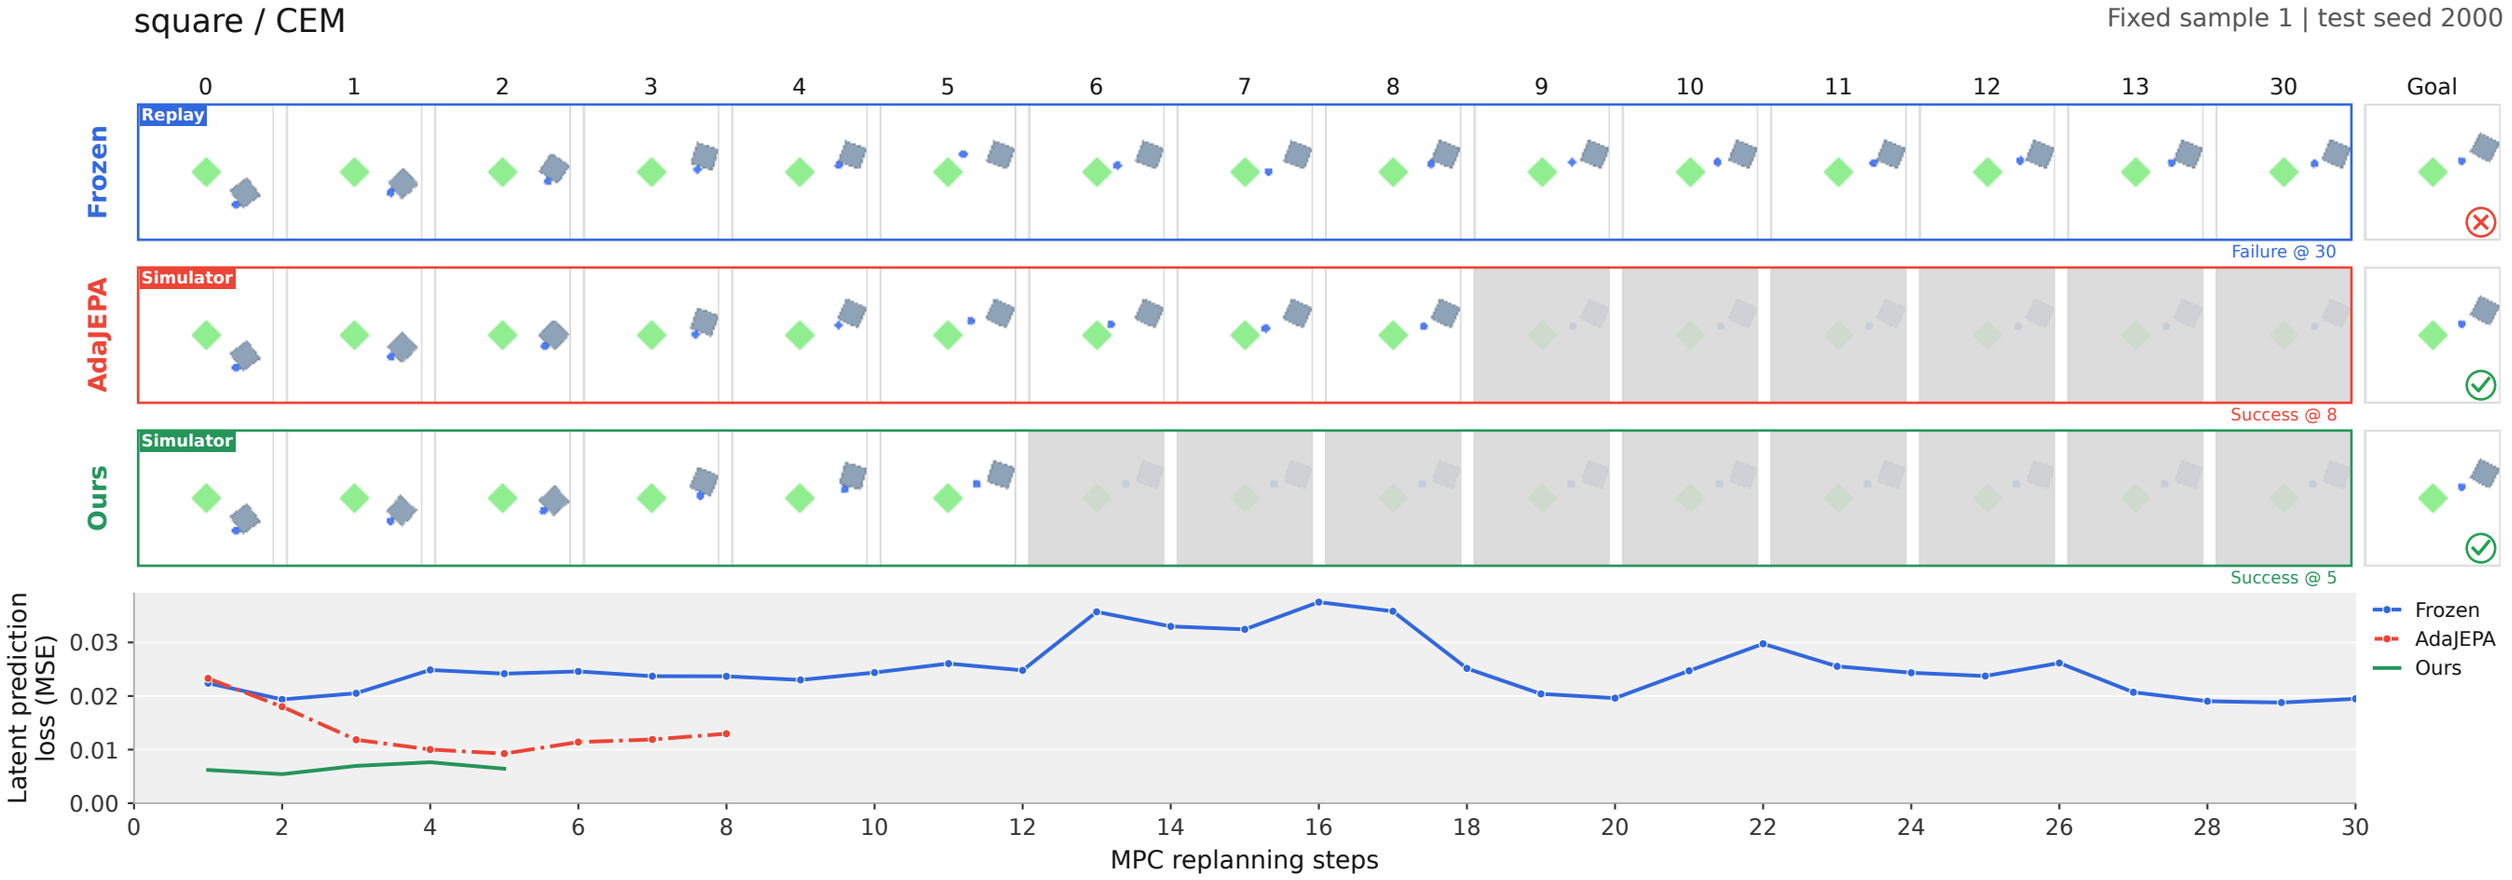


Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.

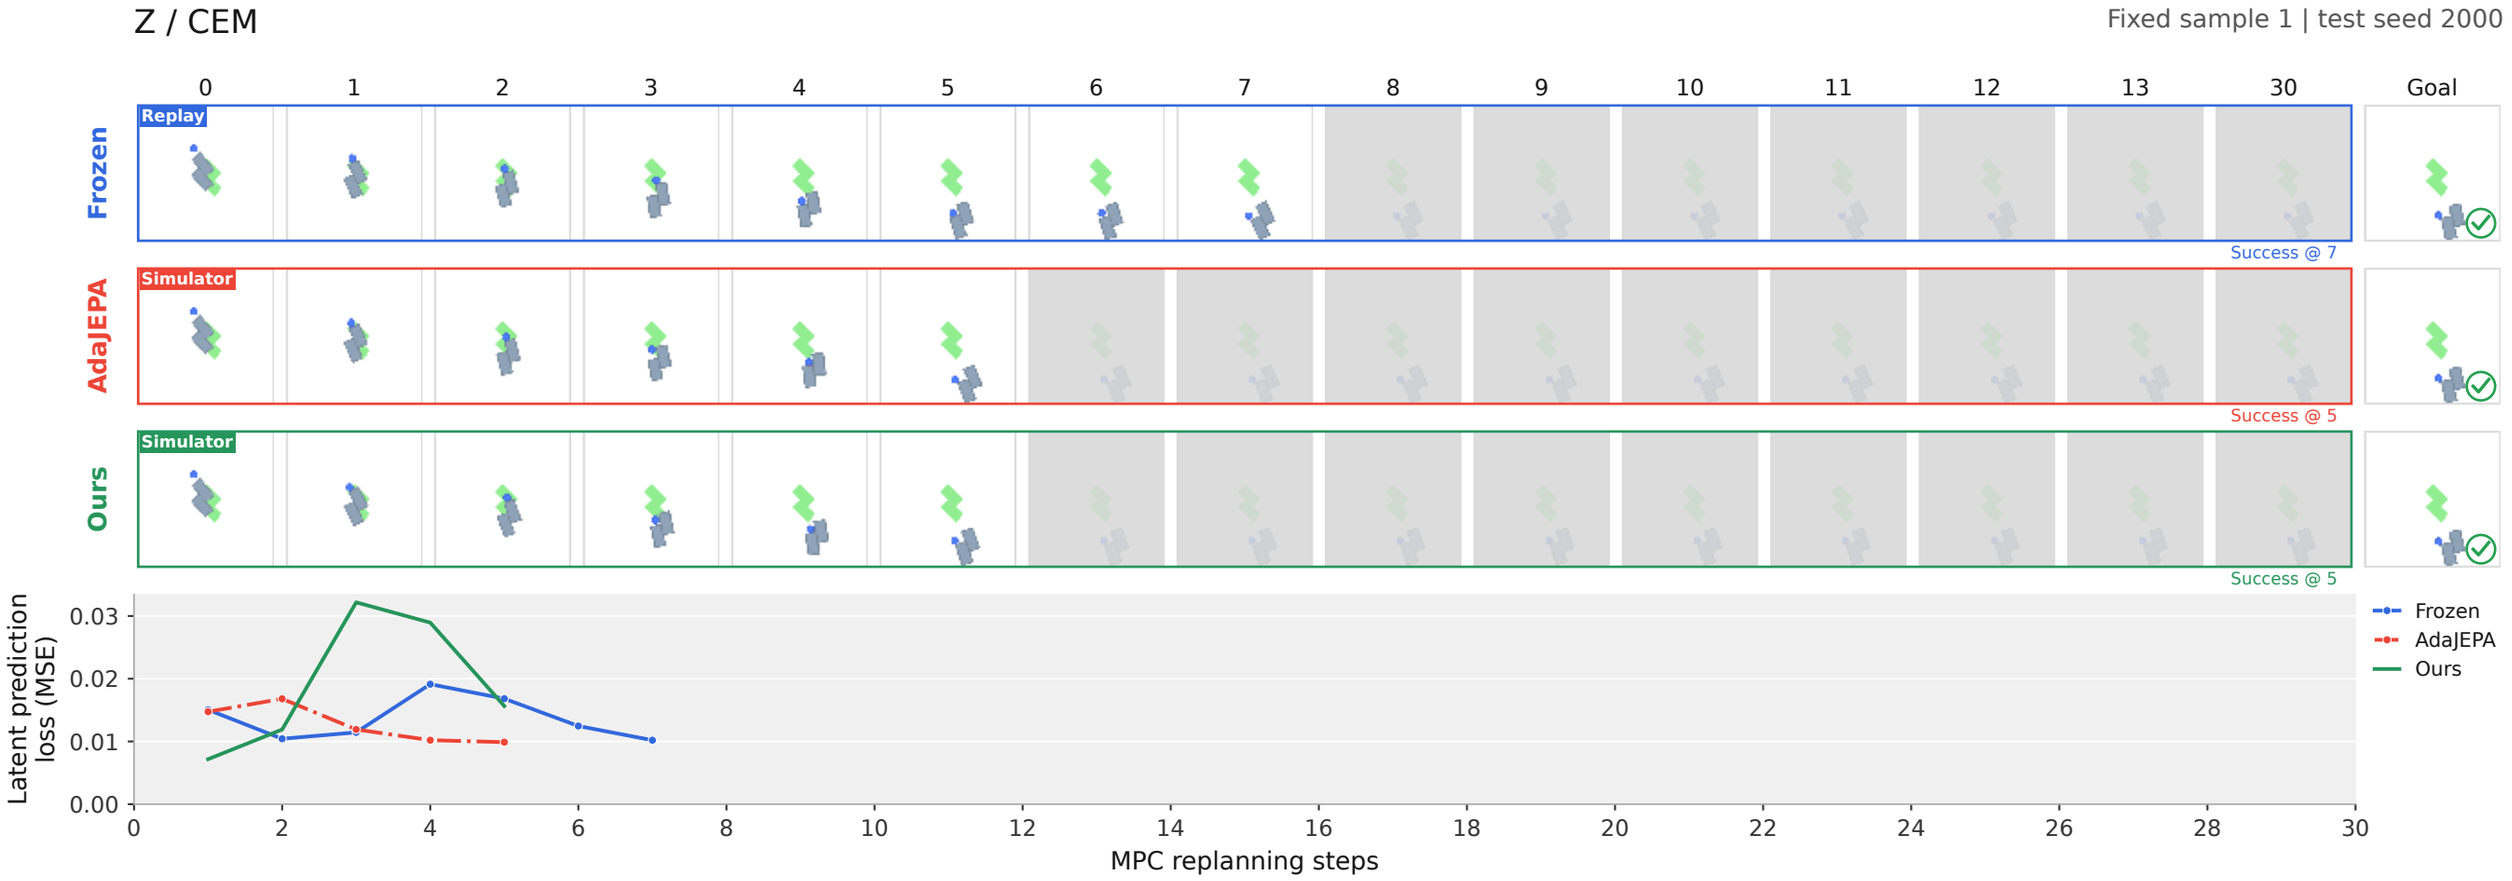


Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.

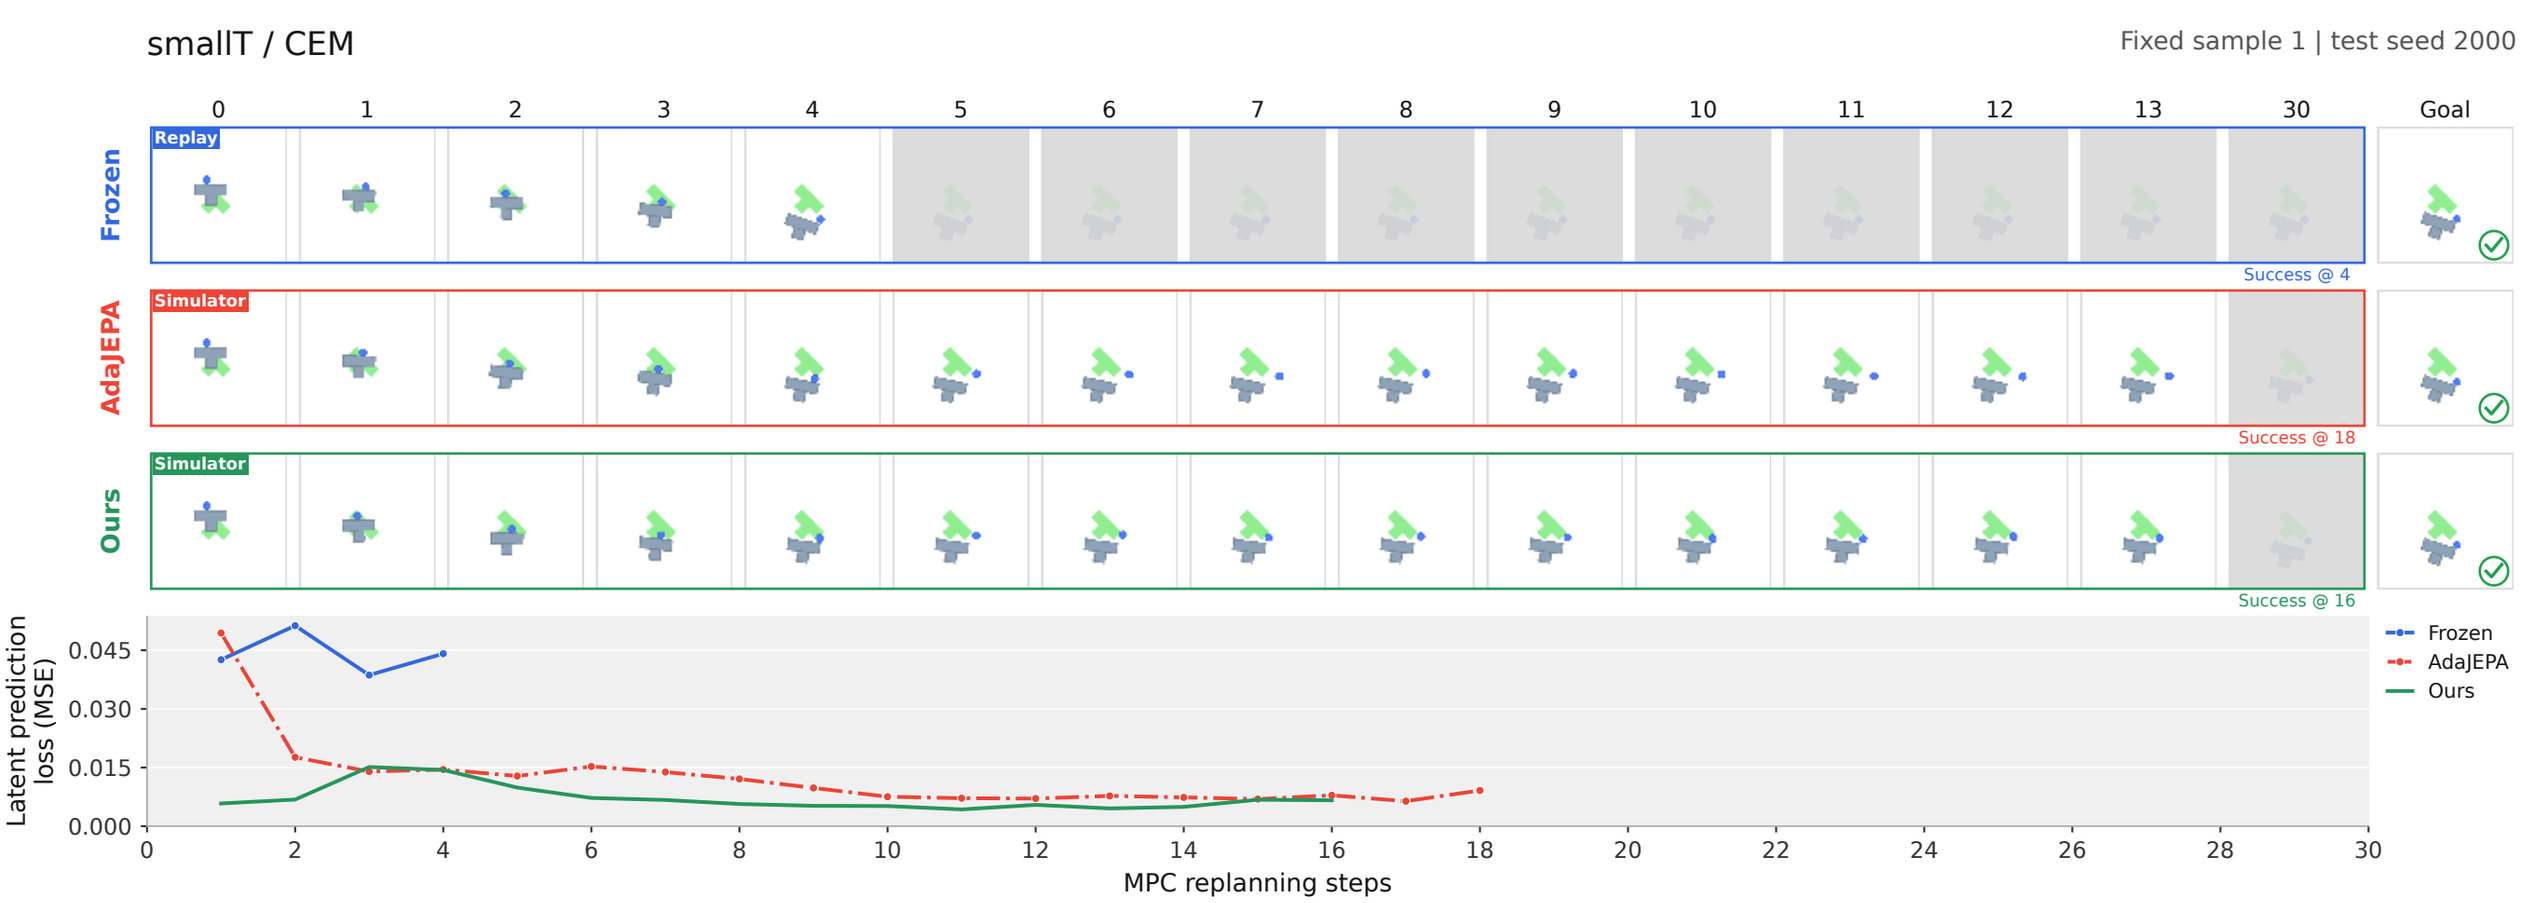
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.



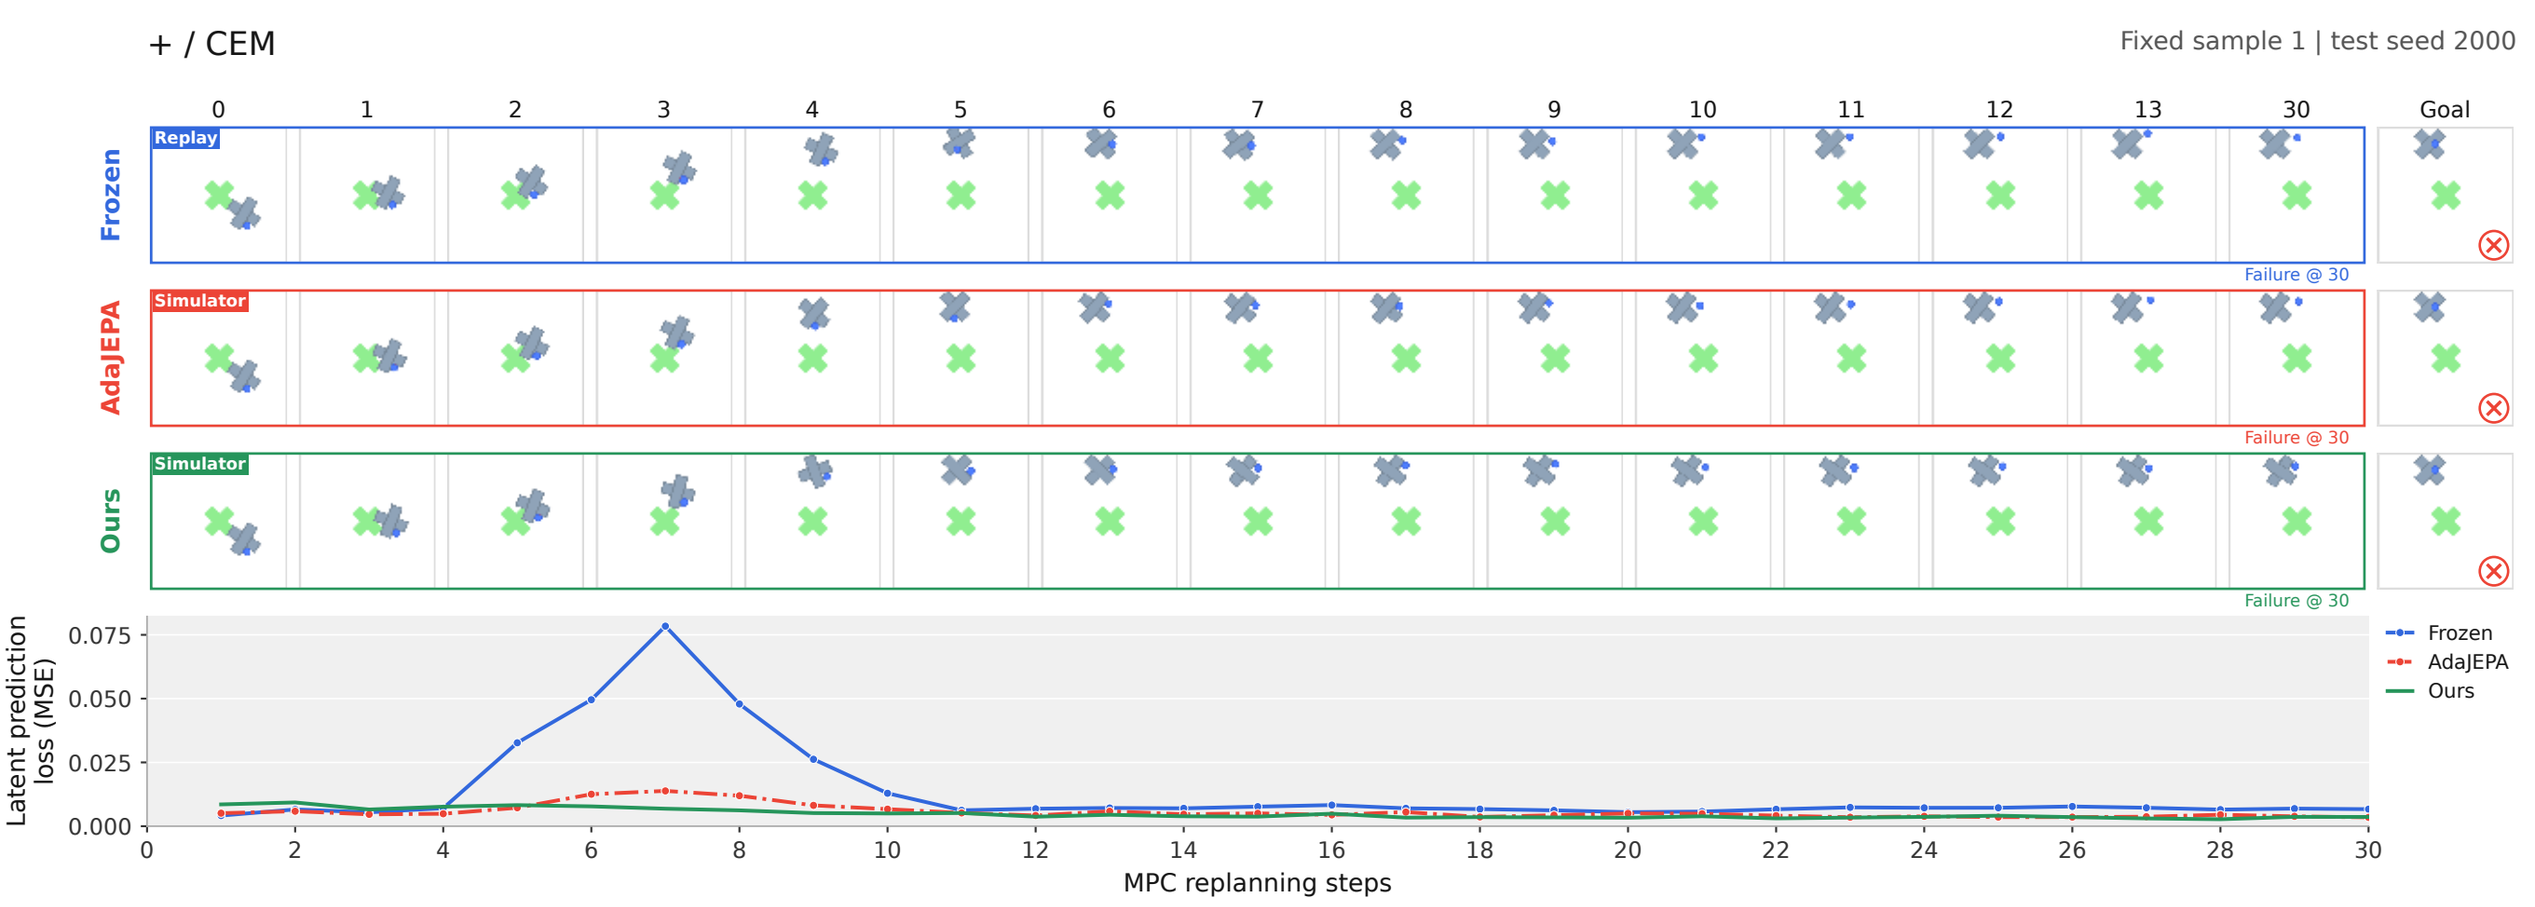
Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.



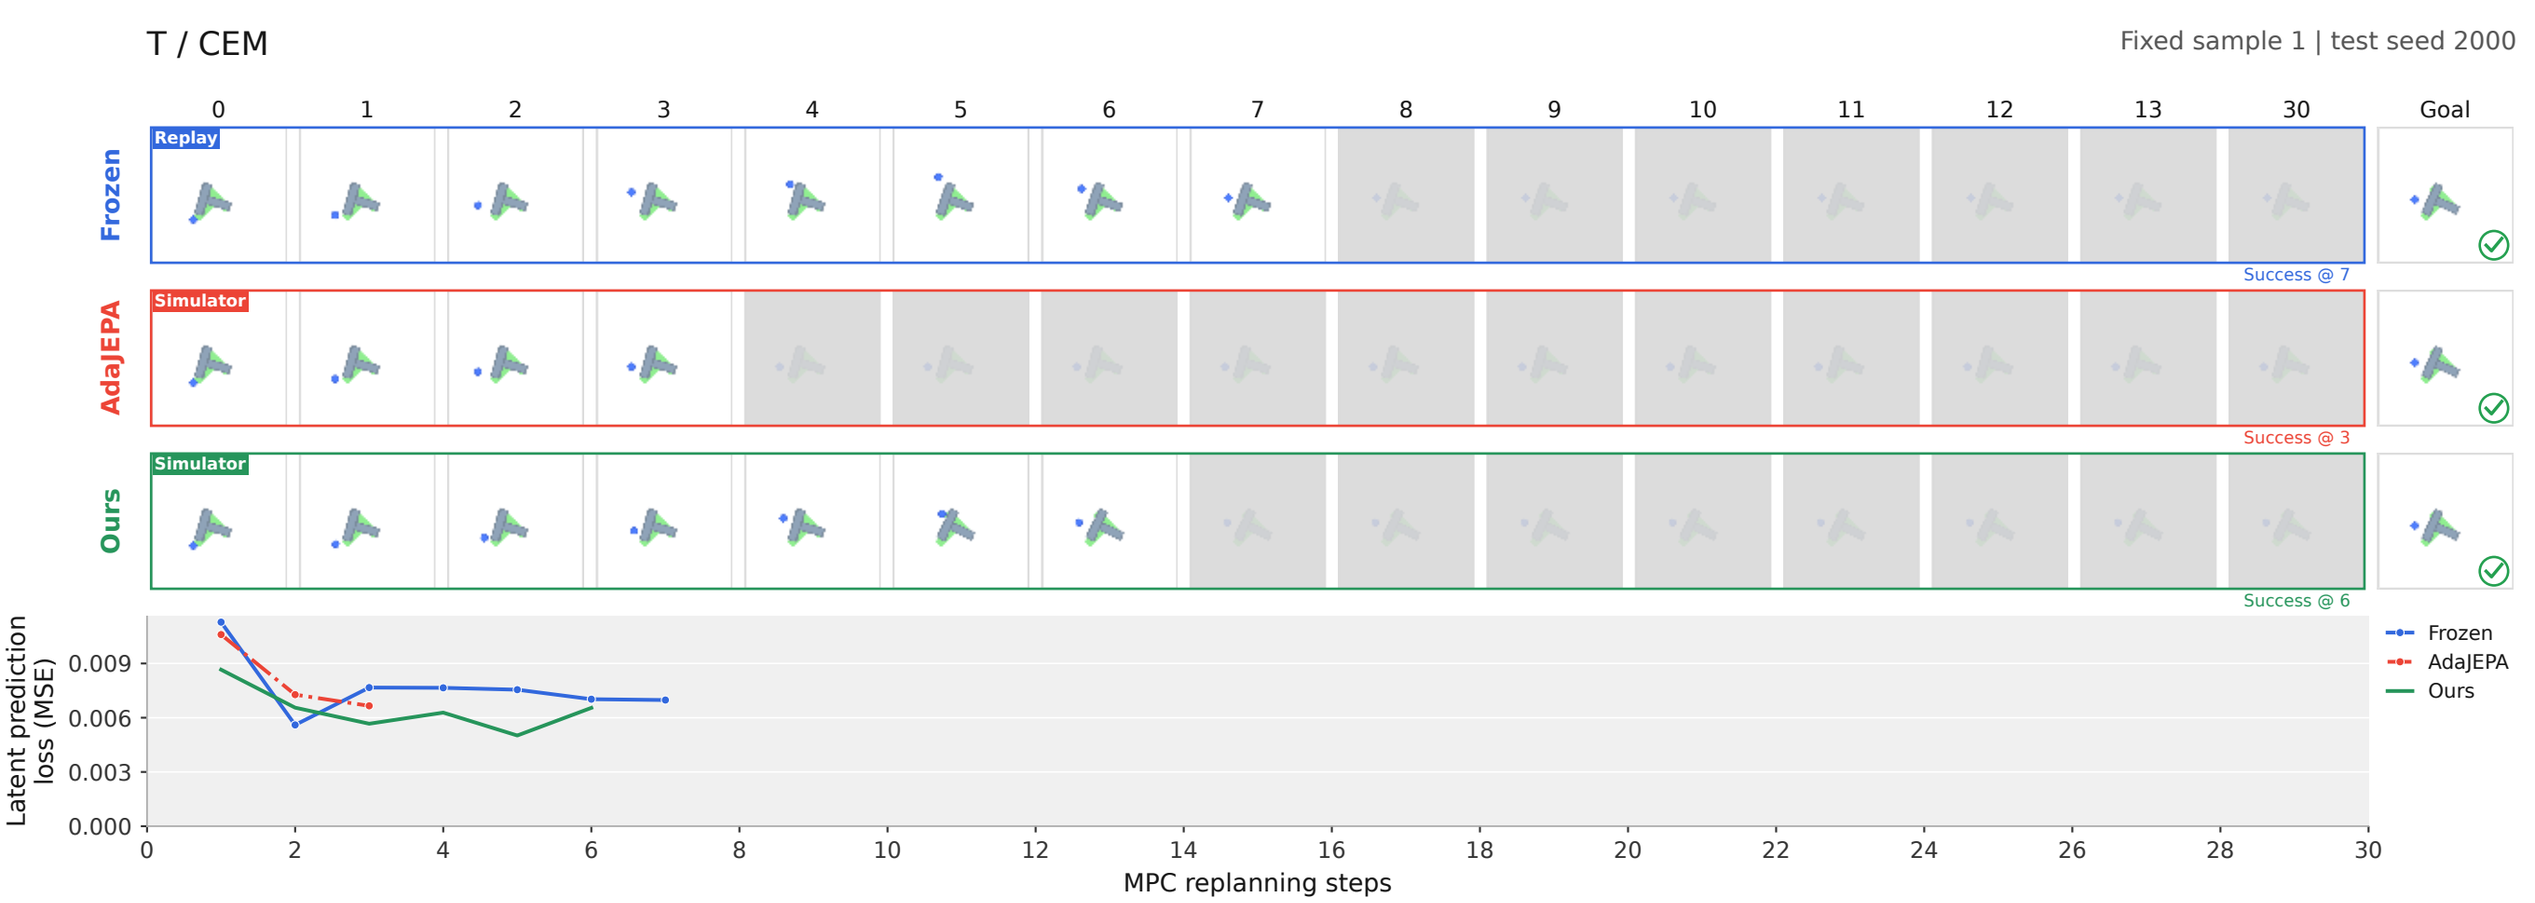
Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.



Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.



Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.



Prediction MSE on up to five recent transitions. Visual + proprioceptive latents; action dimensions excluded. Frozen: same criterion, computed on its replayed trajectory.
Gray frames hold the terminal state. Curves stop at termination. Each method is scored with its own weights, so raw loss levels are not directly comparable.

MPC planning steps - held-out shape-shift targets

Seven shapes x 50 targets x 3 test seeds per method and planner; identical initial/goal pairs. Ours: round-5 states of three chains, read-only.

Planner	Method	Episodes	Success (%)	Steps, all episodes	Steps, successes only
GD	Frozen	1,050	42.48	20.02 +/- 0.37	6.52 +/- 0.10
GD	AdaJEPA	1,050	72.19	14.52 +/- 0.61	8.56 +/- 0.11
GD	**Ours**	**3,150**	**81.59**	**11.30 +/- 0.53**	**7.08 +/- 0.15**
CEM	Frozen	1,050	43.81	21.08 +/- 0.58	9.65 +/- 0.16
CEM	AdaJEPA	1,050	70.38	16.02 +/- 0.33	10.13 +/- 0.45
CEM	**Ours**	**3,150**	**79.08**	**13.53 +/- 0.29**	**9.17 +/- 0.02**

Episodes per method and planner: 7 shapes x 50 targets x 3 test seeds = 1,050 (Ours: 3,150, its three chains on the same 1,050 targets).
Steps are MPC replanning iterations. Mean +/- sample SD is across the three test-seed means; lower means fewer iterations.
All-episode means include failures, which each used the full 30-step budget. Success-only means use each method's successful subset.
Success-only subsets differ across methods; read both step columns with the success rate. Checkpoints differ between Ours and the baselines.
Source: final_eval_results.npz of every run (action_len = first successful replan), checked against per_iter_logs.json for the sequential runs.